

Predicting Daily Bike-Sharing Demand (UCI Capital Bikeshare) - Regression

Design and Creative Technologies, Torrens University

- **Student:** Luis Guilherme de Barros Andrade Faria - A00187785
- **Subject Name:** Machine Learning
- **Subject Code:** MLN 601
- **Title:** Predicting Daily Bike Rental Demand - multi-model regression
- **Lecturer:** Dr. Kamran Shaukat
- **Assessment No.:** 3
- **Date:** August 2026

Field	Value
Dataset	UCI Bike Sharing - daily counts, Capital Bikeshare D.C. (Fanaee-T & Gama, 2014)
Task	Regression: predict the daily total rental count (cnt)
Algorithms	Linear Regression + KNN + Random Forest + Gradient Boosting (vs a mean baseline)
Tuning	GridSearchCV, 5-fold cross-validation, scoring = negative MAE
Evaluation	MAE (primary), RMSE, R-squared
Methodology	CRISP-DM

This notebook follows the six CRISP-DM stages to predict the **daily** number of bikes rented from the Capital Bikeshare system in Washington D.C., using calendar and weather features. It compares four learning algorithms against a mean baseline and tunes them with cross-validation, reporting the **mean absolute error** (the metric the brief asks for) alongside RMSE and R-squared.

Demand is an operational planning signal, so the models are framed as a **decision-support tool** for fleet rebalancing and staffing, not an exact forecast.

1. Business Understanding

Bike-sharing systems such as Capital Bikeshare let users rent a bike at one station and return it at another. Operators face a daily logistics problem: too few bikes at a station means lost trips and frustrated riders, while too many means wasted redistribution effort. Being able to **predict how many bikes will be rented on a given day** - from the

calendar and the weather forecast – directly supports fleet rebalancing, staffing, and maintenance planning, and helps riders find a bike when they need one.

The **business question** is therefore: *can we predict the total daily rental count from routine calendar and weather information?* Framed in machine-learning terms, this is a **supervised regression** task: the target is `cnt` (the total number of daily rentals), and the inputs are season, year, month, weekday, holiday/working-day flags, the weather situation, and normalised temperature, humidity and wind speed.

Two assumptions shape the work. First, the relationship between weather/calendar and demand is **non-linear** (demand rises with temperature up to a comfortable point, then can fall), so tree-based ensembles are expected to beat a straight line. Second, because demand grew strongly as the system matured, **year** is an important signal. Success is judged against a naive baseline that always predicts the mean: a useful model must beat it and achieve a low **MAE**, which is directly interpretable as "average error in bikes per day". For the operator, even a model that is roughly right is valuable: a day-ahead demand estimate lets staff pre-position bikes at busy stations, schedule maintenance on quiet days, and size the casual-rider supply for warm weekends, turning a reactive logistics problem into a planned one.

2. Data Understanding

The data is the UCI **Bike Sharing** dataset (Fanaee-T & Gama, 2014), specifically the **daily** file `day.csv` – 731 days spanning 2011-2012 for Capital Bikeshare in Washington D.C., combining rental counts with weather and seasonal information. Daily (rather than hourly) aggregation is used deliberately: it answers the day-ahead planning question directly and avoids the strong hour-of-day cycle dominating the analysis. The dependent variable is `cnt`.

```
In [1]: import os
import warnings

# Silence sklearn/joblib parallel UserWarnings so they cannot flood the export
# PDF. n_jobs=1 in the grid searches below means no worker processes spawned
os.environ["PYTHONWARNINGS"] = "ignore"
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.dummy import DummyRegressor
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.preprocessing import StandardScaler, OneHotEncoder
```

```

from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

sns.set_theme(style="whitegrid")
RANDOM_STATE = 42

```

```

In [2]: from pathlib import Path

NB_DIR = Path.cwd()
BASE_DIR = NB_DIR.parent if NB_DIR.name == "notebook" else NB_DIR
OUTPUT_DIR = BASE_DIR / "outputs"
FIG_DIR = OUTPUT_DIR / "figures"
FIG_DIR.mkdir(parents=True, exist_ok=True)

UCI_URL = ("https://archive.ics.uci.edu/static/public/275/"
          "bike+sharing+dataset.zip")

def load_bike():
    candidates = [BASE_DIR / "dataset" / "day.csv", NB_DIR / "dataset" / "day.csv",
                  NB_DIR / "day.csv"]
    local = next((p for p in candidates if p.exists()), None)
    if local is not None:
        print("Loading local:", local)
        return pd.read_csv(local, parse_dates=["dteday"])
    import io, zipfile, urllib.request
    print("Downloading from UCI:", UCI_URL)
    with urllib.request.urlopen(UCI_URL) as r:
        zf = zipfile.ZipFile(io.BytesIO(r.read()))
    return pd.read_csv(zf.open("day.csv"), parse_dates=["dteday"])

df = load_bike()
print("Shape:", df.shape, "| date range:", df.dteday.min().date(), "->", df.dteday.max().date())
df.head()

```

Loading local: /Users/luisfaria/Desktop/sEngineer/masters_SWEAI/2026-T2/MLN/assignments/Assessment3/dataset/day.csv
Shape: (731, 16) | date range: 2011-01-01 -> 2012-12-31

```
Out[2]:
```

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	weathersit
0	1	2011-01-01	1	0	1	0	6	0	2 0.3
1	2	2011-01-02	1	0	1	0	0	0	2 0.3
2	3	2011-01-03	1	0	1	0	1	1	1 0.1
3	4	2011-01-04	1	0	1	0	2	1	1 0.2
4	5	2011-01-05	1	0	1	0	3	1	1 0.2

```
In [3]: df.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 731 entries, 0 to 730
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   instant     731 non-null    int64
1   dteday      731 non-null    datetime64[us]
2   season      731 non-null    int64
3   yr          731 non-null    int64
4   mnth        731 non-null    int64
5   holiday     731 non-null    int64
6   weekday     731 non-null    int64
7   workingday  731 non-null    int64
8   weathersit   731 non-null    int64
9   temp        731 non-null    float64
10  atemp       731 non-null    float64
11  hum         731 non-null    float64
12  windspeed   731 non-null    float64
13  casual      731 non-null    int64
14  registered  731 non-null    int64
15  cnt         731 non-null    int64
dtypes: datetime64[us](1), float64(4), int64(11)
memory usage: 91.5 KB
```

```
In [4]: df.describe().T
```

Out [4]:

	count	mean	min	25%	50%	75%	max
instant	731.0	366.0	1.0	183.5	366.0	548.5	731.0
dteday	731	2012-01-01 00:00:00	2011-01-01 00:00:00	2011-07-02 12:00:00	2012-01-01 00:00:00	2012-07-01 12:00:00	2012-12-31 00:00:00
season	731.0	2.49658	1.0	2.0	3.0	3.0	4.0
yr	731.0	0.500684	0.0	0.0	1.0	1.0	1.0
mnth	731.0	6.519836	1.0	4.0	7.0	10.0	12.0
holiday	731.0	0.028728	0.0	0.0	0.0	0.0	1.0
weekday	731.0	2.997264	0.0	1.0	3.0	5.0	6.0
workingday	731.0	0.683995	0.0	0.0	1.0	1.0	1.0
weathersit	731.0	1.395349	1.0	1.0	1.0	2.0	3.0
temp	731.0	0.495385	0.05913	0.337083	0.498333	0.655417	0.861667
atemp	731.0	0.474354	0.07907	0.337842	0.486733	0.608602	0.840896
hum	731.0	0.627894	0.0	0.52	0.626667	0.730209	0.9725
windspeed	731.0	0.190486	0.022392	0.13495	0.180975	0.233214	0.507463
casual	731.0	848.176471	2.0	315.5	713.0	1096.0	3410.0
registered	731.0	3656.172367	20.0	2497.0	3662.0	4776.5	6946.0
cnt	731.0	4504.348837	22.0	3152.0	4548.0	5956.0	8714.0

```
In [5]: print("Missing values per column:")
print(df.isna().sum().to_string())
print("\nDuplicate rows:", int(df.duplicated().sum()))
print("Identity check (casual + registered == cnt):",
      bool((df["casual"] + df["registered"] == df["cnt"]).all()))
```

Missing values per column:

instant	0
dteday	0
season	0
yr	0
mnth	0
holiday	0
weekday	0
workingday	0
weathersit	0
temp	0
atemp	0
hum	0
windspeed	0
casual	0
registered	0
cnt	0

Duplicate rows: 0

Identity check (casual + registered == cnt): True

2.2 Variables

The dataset has 16 columns. The **target** is `cnt`, the total daily rentals. Helpful structure:

- **Calendar:** `season` (1 spring, 2 summer, 3 fall, 4 winter), `yr` (0 = 2011, 1 = 2012), `mnth` (1-12), `weekday` (0-6), `holiday`, `workingday`.
- **Weather:** `weathersit` (1 clear, 2 mist, 3 light rain/snow, 4 heavy rain/snow) and the **normalised** measures `temp` (/41 C), `atemp` (feels-like /50 C), `hum` (/100) and `windspeed` (/67).
- **Counts:** `casual` and `registered` users, whose sum is exactly `cnt` (verified above).

Because `casual + registered = cnt` by construction, **both are dropped from the model inputs** - using them would be target leakage. `instant` (a row index) and the raw `dteday` are also dropped, since their information already lives in the calendar columns.

2.3 Exploratory data analysis

The plots below explore demand against time, weather and the calendar, following the variables the brief highlights: season / month / weekday against `cnt`, weather against `cnt`, the continuous weather measures against `cnt`, and the casual-versus-registered split. Helper label columns are added for readability (they are not used as model inputs).


```
In [6]: season_lbl = {1: "spring", 2: "summer", 3: "fall", 4: "winter"}
        weather_lbl = {1: "clear", 2: "mist", 3: "light rain/snow", 4: "heavy rain/s
```

```

weekday_lbl = {0: "Sun", 1: "Mon", 2: "Tue", 3: "Wed", 4: "Thu", 5: "Fri", 6: "Sat"}
df["season_name"] = df["season"].map(season_lbl)
df["weather_name"] = df["weathersit"].map(weather_lbl)
df["weekday_name"] = df["weekday"].map(weekday_lbl)
df["year_name"] = df["yr"].map({0: "2011", 1: "2012"})

# Distribution of the target and its growth over the two years
fig, ax = plt.subplots(1, 2, figsize=(13, 4))
sns.histplot(df["cnt"], bins=30, kde=True, ax=ax[0])
ax[0].set_title("Distribution of daily rentals (cnt)")
ax[0].set_xlabel("daily rental count")
sns.lineplot(data=df, x="dteday", y="cnt", hue="year_name", ax=ax[1])
ax[1].set_title("Daily rentals over time (2011-2012)")
ax[1].set_xlabel("date"); ax[1].set_ylabel("cnt")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_cnt_distribution.png", dpi=120)
plt.show()

```




Read. The target is right-skewed around a typical ~4,500 rentals a day, and the time series shows a clear upward step from 2011 to 2012 on top of a seasonal wave - so both `yr` and the seasonal calendar will carry signal.

```


In [7]: # Season / month / weekday against total rentals (rubric-named plots)
fig, ax = plt.subplots(1, 3, figsize=(16, 4))
sns.barplot(data=df, x="season_name", y="cnt",
            order=["spring", "summer", "fall", "winter"], ax=ax[0])
ax[0].set_title("Mean rentals by season"); ax[0].set_xlabel("")
sns.lineplot(data=df, x="mnth", y="cnt", hue="year_name", marker="o", ax=ax[1])
ax[1].set_title("Mean rentals by month"); ax[1].set_xlabel("month")
sns.barplot(data=df, x="weekday_name", y="cnt",
            order=["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"], ax=ax[2])
ax[2].set_title("Mean rentals by weekday"); ax[2].set_xlabel("")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_season_month_weekday.png", dpi=120)
plt.show()

```




Read. Demand peaks in summer and fall and bottoms in spring/winter; every month sits higher in 2012 than 2011; and weekdays differ little - the weekly cycle is weak for the daily total, unlike the strong seasonal one.

```
In [8]: # Weather situation and working-day effect on rentals
fig, ax = plt.subplots(1, 2, figsize=(13, 4))
sns.boxplot(data=df, x="weather_name", y="cnt",
            order=["clear", "mist", "light rain/snow", "heavy rain/snow"], ax=ax[0])
ax[0].set_title("Rentals by weather situation"); ax[0].set_xlabel("")
sns.boxplot(data=df, x="workingday", y="cnt", ax=ax[1])
ax[1].set_title("Rentals: working day (1) vs weekend/holiday (0)")
ax[1].set_xlabel("working day")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_weather_cnt.png", dpi=120)
plt.show()
```




Read. Clear days carry the most rentals and demand falls step-wise as the weather worsens, with light or heavy rain/snow the lowest. Working days and weekends share a similar median, as commuter and leisure demand offset each other.

```
In [9]: # Continuous weather measures against rentals
cont = ["temp", "atemp", "hum", "windspeed"]
fig, ax = plt.subplots(1, 4, figsize=(17, 4))
for a, c in zip(ax, cont):
    sns.regplot(data=df, x=c, y="cnt", scatter_kws={"alpha": 0.3, "s": 12},
                line_kws={"color": "black"}, ax=a)
    a.set_title(f"{c} vs cnt")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_continuous_vs_cnt.png", dpi=120)
plt.show()
```

Read. Temperature and feels-like temperature (`atemp`) rise strongly with demand - the clearest positive drivers - while humidity and especially wind speed trend mildly negative.


```
In [10]: # Casual vs registered composition of total demand
comp = df.groupby("season_name")[["casual", "registered"]].mean().reindex(
    ["spring", "summer", "fall", "winter"])
fig, ax = plt.subplots(1, 2, figsize=(13, 4))
comp.plot(kind="bar", stacked=True, ax=ax[0])
ax[0].set_title("Mean casual vs registered by season")
ax[0].set_xlabel(""); ax[0].set_ylabel("mean daily users")
ax[0].tick_params(axis="x", rotation=0)
ax[1].stackplot(df["dteday"], df["casual"], df["registered"],
    labels=["casual", "registered"], alpha=0.8)
ax[1].set_title("Casual + registered over time = cnt")
ax[1].set_xlabel("date"); ax[1].legend(loc="upper left")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_casual_registered.png", dpi=120)
plt.show()
```



Read. Registered commuters make up the bulk of demand and stay relatively steady year-round, while casual riders swell in the warm seasons - the two segments respond differently to weather, which the single `cnt` target blends together.

```
In [11]: # Correlation heatmap of the modelling variables (+ cnt)
num_cols = ["season", "yr", "mnth", "holiday", "weekday", "workingday",
    "weathersit", "temp", "atemp", "hum", "windspeed", "cnt"]
corr = df[num_cols].corr()
plt.figure(figsize=(10, 8))
sns.heatmap(corr, annot=True, fmt=".2f", cmap="coolwarm", center=0,
    square=True, cbar_kws={"shrink": 0.8}, annot_kws={"size": 7})
plt.title("Correlation heatmap (features + cnt)")
plt.tight_layout()
```

```
plt.savefig(FIG_DIR / "v1_corr_heatmap.png", dpi=120)
plt.show()
corr["cnt"].drop("cnt").sort_values(ascending=False).round(3).to_frame("corr")
```




Out [11]:

	corr_with_cnt
atemp	0.631
temp	0.627
yr	0.567
season	0.406
mnth	0.280
weekday	0.067
workingday	0.061
holiday	-0.068
hum	-0.101
windspeed	-0.235
weathersit	-0.297

Read. The heatmap confirms the EDA numerically: `atemp` and `temp` (~0.63) and `yr` (0.57) are the strongest positive correlations with `cnt`, while `weathersit` (-0.30) and `windspeed` (-0.24) are the main negatives. `temp` and `atemp` are near-collinear (~0.99), which destabilises a linear model's coefficients but not the trees.

```
In [12]: # Seaborn pairplot (required by the brief) - key continuous drivers + target
g = sns.pairplot(df[["temp", "atemp", "hum", "windspeed", "cnt"]],
                 corner=True, plot_kws={"alpha": 0.3, "s": 12}, diag_kind="kde")
g.fig.suptitle("Pairplot of continuous weather variables and cnt", y=1.02)
g.savefig(FIG_DIR / "v1_pairplot.png", dpi=110)
plt.show()
```

Pairplot of continuous weather variables and cnt



EDA summary. Daily demand is driven mostly by **temperature and season**, lifted by the **2011 -> 2012 growth**, and suppressed by **adverse weather, humidity and wind**. The weekly cycle is weak for the daily total because steady registered commuters dominate and offset the more weather-sensitive casual riders. There are no missing values, and the relationships look distinctly non-linear - which should favour the tree ensembles over a straight line.

3. Data Preparation

Preparation has three parts. **(1) Feature selection / leakage control:** `casual`, `registered` (which sum to the target), `instant` (an index) and the raw `dteday` and helper label columns are dropped, leaving eleven predictors. **(2) Encoding and scaling, inside a Pipeline:** the nominal calendar/weather codes (`season` , `mnth` ,

`weekday` , `weathersit`) are **one-hot encoded** so linear and distance models do not read them as ordered numbers; the continuous measures (`temp` , `atemp` , `hum` , `windspeed`) are **standard-scaled** (needed by KNN and linear regression, harmless for trees); the binary flags (`yr` , `holiday` , `workingday`) pass through. Wrapping this `ColumnTransformer` **inside each model's Pipeline** means it is refit on each training fold during cross-validation and never sees the validation rows - preventing data leakage. **(3) Train/test split:** a 75/25 split with a fixed seed, as the brief suggests. The two years are pooled and `yr` is kept as a feature, so the models can use both 2011 and 2012; a strict time-ordered (2011 -> 2012) split is noted as a stronger forecasting test for future work.

```
In [13]: DROP = ["casual", "registered", "cnt", "instant", "dteday",
                "season_name", "weather_name", "weekday_name", "year_name"]
nominal = ["season", "mnth", "weekday", "weathersit"]
binary = ["yr", "holiday", "workingday"]
continuous = ["temp", "atemp", "hum", "windspeed"]
feature_cols = nominal + binary + continuous

X = df[feature_cols]
y = df["cnt"]

preprocessor = ColumnTransformer([
    ("oh", OneHotEncoder(drop="first", handle_unknown="ignore"), nominal),
    ("sc", StandardScaler(), continuous),
    ("pass", "passthrough", binary),
])

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.25, random_state=RANDOM_STATE)

print("Train:", X_train.shape, "| Test:", X_test.shape)
print("\nTop correlations with cnt:")
print(df[feature_cols + ["cnt"]].corr()["cnt"].drop("cnt")
      .sort_values(ascending=False).round(3).to_string())
```

Train: (548, 11) | Test: (183, 11)

Top correlations with cnt:

atemp	0.631
temp	0.627
yr	0.567
season	0.406
mnth	0.280
weekday	0.067
workingday	0.061
holiday	-0.068
hum	-0.101
windspeed	-0.235
weathersit	-0.297

4. Modelling

Five estimators are compared so the ensembles can be read in context:

1. **Mean baseline** (`DummyRegressor`) - always predicts the training mean; the floor every real model must beat.
2. **Linear Regression** - an interpretable straight-line baseline.
3. **K-Nearest Neighbors** - a simple non-linear, distance-based learner; `n_neighbors` tuned.
4. **Random Forest** - a bagging **ensemble** of trees; `n_estimators` , `max_depth` and `min_samples_leaf` tuned.
5. **Gradient Boosting** - a boosting **ensemble**; `n_estimators` , `max_depth` and `learning_rate` tuned.

Going beyond two models, and including two different ensembles, lets us see whether bagging or boosting suits this data better. Each model is a `Pipeline(preprocessor -> estimator)` , and the tuned models use `GridSearchCV` with **5-fold cross-validation** scored on **negative MAE** (the brief's metric). All estimators use a fixed seed and `n_jobs=1` for deterministic, warning-free runs; the held-out test set is touched once, in Evaluation.

```
In [14]: def pipe(est):
          return Pipeline([("prep", preprocessor), ("model", est)])

searches = {
    "K-Nearest Neighbors": (KNeighborsRegressor(), {
        "model__n_neighbors": [3, 5, 7, 9, 11, 15],
        "model__weights": ["uniform", "distance"]}),
    "Random Forest": (RandomForestRegressor(random_state=RANDOM_STATE), {
        "model__n_estimators": [200, 400],
        "model__max_depth": [None, 8, 16],
        "model__min_samples_leaf": [1, 2, 4]}),
    "Gradient Boosting": (GradientBoostingRegressor(random_state=RANDOM_STATE), {
        "model__n_estimators": [200, 400],
        "model__max_depth": [2, 3],
        "model__learning_rate": [0.05, 0.1]}),
}

fitted = {}
fitted["Mean baseline"] = pipe(DummyRegressor(strategy="mean")).fit(X_train, y_train)
fitted["Linear Regression"] = pipe(LinearRegression()).fit(X_train, y_train)

best_params = {}
for name, (est, grid) in searches.items():
    gs = GridSearchCV(pipe(est), grid, scoring="neg_mean_absolute_error",
                      cv=5, n_jobs=1)
    gs.fit(X_train, y_train)
    fitted[name] = gs.best_estimator_
    best_params[name] = gs.best_params_
    print(f"{name}: CV MAE = {-gs.best_score_: .1f} | best = {gs.best_params_}")
```

K-Nearest Neighbors: CV MAE = 752.9 | best = {'model__n_neighbors': 5, 'model__weights': 'distance'}

Random Forest: CV MAE = 538.9 | best = {'model__max_depth': 16, 'model__min_samples_leaf': 1, 'model__n_estimators': 400}
Gradient Boosting: CV MAE = 509.8 | best = {'model__learning_rate': 0.05, 'model__max_depth': 2, 'model__n_estimators': 400}

4.2 Selected hyperparameters

The tuned configurations below were refit on the full training set. They are saved to `outputs/` for traceability.

```
In [15]: bp = pd.DataFrame({k: {p.replace("model__", ""): v for p, v in params.items()
                                for k, params in best_params.items()}).T
bp.to_csv(OUTPUT_DIR / "rf_best_params_v1.csv")
bp
```

```
Out[15]:
```

	n_neighbors	weights	max_depth	min_samples_leaf	n_estimators	learning_rate
K-Nearest Neighbors	5	distance	NaN	NaN	NaN	NaN
Random Forest	NaN	NaN	16.0	1.0	400.0	NaN
Gradient Boosting	NaN	NaN	2.0	NaN	400.0	0.05

5. Evaluation

The fitted models are judged on the **held-out 25% test set**. Three complementary regression metrics are reported:

- **MAE** (mean absolute error) - the brief's metric, directly readable as the average error in **bikes per day**.
- **RMSE** (root mean squared error) - penalises large misses more heavily, so it flags days the model gets badly wrong.
- **R-squared** - the share of the variance in daily demand the model explains, relative to the mean baseline.

A predicted-versus-actual plot, a residual plot and a feature-importance chart then make the best model's behaviour and logic transparent.

```
In [16]: def reg_metrics(model, X_te, y_te):
pred = model.predict(X_te)
mse = mean_squared_error(y_te, pred)
return {"MAE": mean_absolute_error(y_te, pred),
        "RMSE": np.sqrt(mse),
        "R2": r2_score(y_te, pred)}

results = pd.DataFrame({name: reg_metrics(m, X_test, y_test)
                        for name, m in fitted.items()}).T.round(3)
```

```

results.index.name = "model"
results = results.sort_values("MAE")
results.to_csv(OUTPUT_DIR / "model_metrics_v1.csv")
results

```

Out[16]:

	MAE	RMSE	R2
model			
Random Forest	450.908	683.297	0.879
Gradient Boosting	473.307	663.210	0.886
Linear Regression	566.091	774.990	0.844
K-Nearest Neighbors	721.113	926.857	0.777
Mean baseline	1662.978	1983.325	-0.021


```

In [17]: # Best model = lowest test MAE (excluding the baseline)
ranked = results.drop(index="Mean baseline")
best_name = ranked["MAE"].idxmin()
best_model = fitted[best_name]
print("Best model by test MAE:", best_name)

pred = best_model.predict(X_test)
fig, ax = plt.subplots(1, 2, figsize=(13, 5))
ax[0].scatter(y_test, pred, alpha=0.4, s=18)
lims = [min(y_test.min(), pred.min()), max(y_test.max(), pred.max())]
ax[0].plot(lims, lims, "k--", label="perfect prediction")
ax[0].set_xlabel("actual cnt"); ax[0].set_ylabel("predicted cnt")
ax[0].set_title(f"Predicted vs actual - {best_name}"); ax[0].legend()
resid = y_test - pred
ax[1].scatter(pred, resid, alpha=0.4, s=18)
ax[1].axhline(0, color="k", ls="--")
ax[1].set_xlabel("predicted cnt"); ax[1].set_ylabel("residual (actual - pred")
ax[1].set_title("Residuals")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_pred_vs_actual.png", dpi=120)
plt.show()


```

Best model by test MAE: Random Forest




```
In [18]: # Feature importance from the best tree-based ensemble
tree_name = best_name if best_name in ("Random Forest", "Gradient Boosting")
tree_model = fitted[tree_name]
feat_names = tree_model.named_steps["prep"].get_feature_names_out()
imp = (pd.DataFrame({"feature": feat_names,
                    "importance": tree_model.named_steps["model"].feature_importances_
                    }.sort_values("importance", ascending=False).reset_index(drop=True)))
imp.to_csv(OUTPUT_DIR / "feature_importance_v1.csv", index=False)

plt.figure(figsize=(8, 6))
sns.barplot(data=imp.head(12), x="importance", y="feature")
plt.title(f"Top feature importances - {tree_name}")
plt.tight_layout()
plt.savefig(FIG_DIR / "v1_feature_importance.png", dpi=120)
plt.show()
imp.head(10)
```



Out[18]:

	feature	importance
0	sc__temp	0.272641
1	pass__yr	0.267824
2	sc__atemp	0.252044
3	sc__hum	0.064178
4	oh__season_4	0.045151
5	sc__windspeed	0.033959
6	oh__weathersit_3	0.006862
7	pass__workingday	0.006085
8	oh__weathersit_2	0.006081
9	oh__mnth_9	0.004688

Interpretation. The tuned **Random Forest** is the best model, with a test **MAE of about 451 bikes per day**, an RMSE of 683 and an **R-squared of 0.88** - it explains the large majority of the variation in daily demand and crushes the mean baseline (MAE 1663). The predicted-versus-actual points hug the diagonal, with residuals scattered fairly evenly around zero; the biggest misses are a handful of unusual days (extreme weather or holidays). Feature importance is led by **feels-like temperature, temperature and the year**, matching the EDA - temperature and the year/season structure drive demand, while wind and humidity play smaller roles. The model comfortably **meets the evaluation criterion** of beating the baseline with a low, interpretable MAE, though it remains a planning aid rather than an exact forecast. Random Forest and Gradient Boosting are effectively tied - the forest takes the lower MAE (the brief's metric) while boosting edges a slightly better RMSE and R-squared - so either ensemble is a defensible final choice, and both decisively beat the linear and KNN models.

6. Deployment / Lessons Learned

In CRISP-DM, deployment for this assessment means reflection rather than a production release. **What worked:** framing demand as regression and comparing a baseline, a linear model, KNN and two ensembles made the algorithm trade-off explicit, and tuning on MAE gave an operationally meaningful score (bikes per day). The ensembles clearly beat the straight line, confirming the non-linear, interaction-heavy relationship between weather, calendar and demand that the EDA suggested. A clean `ColumnTransformer` pipeline kept encoding and scaling leakage-free.

What was harder: the strong 2011 -> 2012 growth trend means a model trained on pooled years leans on `yr`; the rare extreme-weather days are under-represented, so

the largest residuals sit there. Daily aggregation was a deliberate simplification that trades the rich hour-of-day signal for a cleaner day-ahead question.

What I would do differently next time: evaluate with a **time-ordered split** (train on 2011, test on 2012) to test true forecasting and expose tree-extrapolation limits; try a **log-transformed target** to stabilise variance at high demand; engineer richer features (a heat-index, lagged demand, public-event flags); and bring in the hourly file for intraday planning. The honest lesson echoes the earlier assessments - careful framing, leakage control and honest evaluation matter as much as the choice of algorithm.

Academic Integrity Declaration

I declare that except where I have referenced, the work I am submitting for this assessment task is my own work. I have read and am aware of the Torrens University Australia Academic Integrity Policy and Procedure viewable online at <http://www.torrens.edu.au/policies-and-forms>

I am aware that I need to keep a copy of all submitted material and their drafts, and I will do so accordingly.

Statement of Acknowledgement

I acknowledge that I have used the following AI tool(s) in the creation of this report:

- Anthropic Claude Opus 4.8
- OpenAI ChatGPT Codex 5.5

Both tools were used to assist with understanding ML concepts, structuring the technical pipeline, improving clarity of academic language, and supporting APA 7th referencing conventions.

Prompt examples:

1. "Explain how to frame daily bike-sharing demand as a regression problem under CRISP-DM, and why mean absolute error is a sensible primary metric for an operational planning use case."
2. "How do I build a leakage-safe scikit-learn pipeline with a ColumnTransformer (one-hot for nominal calendar codes, StandardScaler for continuous weather) and compare Linear Regression, KNN, Random Forest and Gradient Boosting with GridSearchCV?"
3. "Why must casual and registered counts be dropped when predicting cnt in the UCI bike sharing dataset, and how should I interpret feature importances from a Gradient

Boosting model?"

I confirm that the use of these tools has been in accordance with the Torrens University Australia Academic Integrity Policy and TUA, Think and MDS's Position Paper on the Use of AI. I confirm that the final output is authored by me and represents my own critical thinking, analysis, and synthesis of sources. I take full responsibility for the final content of this report.

References

Chapman, P., Clinton, J., Kerber, R., Khabaza, T., Reinartz, T., Shearer, C., & Wirth, R. (2000). *CRISP-DM 1.0: Step-by-step data mining guide*. SPSS Inc. <https://www.the-modeling-agency.com/crisp-dm.pdf>

Fanaee-T, H., & Gama, J. (2014). Event labeling combining ensemble detectors and background knowledge. *Progress in Artificial Intelligence*, 2(2-3), 113-127. <https://doi.org/10.1007/s13748-013-0040-3>

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.

Shaheen, S. A., Guzman, S., & Zhang, H. (2010). Bikesharing in Europe, the Americas, and Asia: Past, present, and future. *Transportation Research Record*, 2143(1), 159-167. <https://doi.org/10.3141/2143-20>

University of California, Irvine. (n.d.). *Bike sharing dataset*. UCI Machine Learning Repository. Retrieved August 2026, from <https://archive.ics.uci.edu/dataset/275/bike+sharing+dataset>