

# Predicting Daily Demand for Capital Bikeshare

## Regression Model Comparison and Temporal Robustness Testing

MLN601 Assessment 3 - CRISP-DM

Design and Creative Technologies, Torrens University

- **Student:** Luis Guilherme de Barros Andrade Faria - A00187785
- **Subject:** Machine Learning (MLN601)
- **Lecturer:** Dr. Kamran Shaukat
- **Assessment:** 3
- **Date:** August 2026

| Field            | Scope                                                                                |
|------------------|--------------------------------------------------------------------------------------|
| Dataset          | UCI Bike Sharing, Capital Bikeshare daily and hourly files, Washington DC, 2011-2012 |
| Target           | cnt : system-wide daily rentals                                                      |
| Primary question | Conditional daily-demand estimation from calendar attributes and expected weather    |
| Secondary check  | Rolling day-ahead temporal robustness, trained in 2011 and evaluated in 2012         |
| Method           | Regression comparison within CRISP-DM (Chapman et al., 2000)                         |

# 1. Business Understanding

## 1.1 Business objective

Capital Bikeshare is a docked bicycle-sharing system (Capital Bikeshare, n.d.). This project asks:

**Given calendar attributes and expected weather conditions, how accurately can regression models estimate Capital Bikeshare's system-wide daily rental demand?**

Daily demand is total network rentals in one day. An accurate estimate could support fleet-capacity, staffing and maintenance planning. Station inventory and rebalancing are outside scope because the CSVs lack station identifiers, locations and capacity; station-level data could extend the work.

Lime in Sydney is a contemporary dockless analogy that motivates my interest and makes the assessment closer to my reality. Although its operations differ, it raises a parallel question: how could calendar and weather patterns inform fleet planning in a local system? It motivates the work but is not model evidence.

## 1.2 Data and prediction scope

Each row represents one system-wide day, with `cnt` as total rentals. Calendar attributes are known in advance. Observed weather stands in for the forecast a real user would hold, so deployment would require forecast inputs. The hourly file supports exploration without changing the daily target.

The primary random 75/25 experiment estimates demand for days resembling 2011-2012 using calendar and weather without past demand. The secondary experiment tests transfer through time: selection uses 2011, then counts through day D plus next-day conditions predict D+1 in 2012. Forecasting all of 2012 on 1 January is separate because later counts would be unavailable.

This report therefore runs three comparisons. The first two answer the assessment question; the third tests whether that answer transfers forward in time.

**Table 1 - What this report compares**

| # | Comparison         | Split                                         | Compared against                                   | Question answered                                         | Where          |
|---|--------------------|-----------------------------------------------|----------------------------------------------------|-----------------------------------------------------------|----------------|
| 1 | Model selection    | Random 75/25, training partition only         | The other 11 family and feature-set configurations | Which configuration predicts most accurately?             | Table 6        |
| 2 | Value of modelling | Random 75/25, holdout opened once             | A constant training-mean prediction                | Does the selected model beat not modelling at all?        | Tables 7 and 8 |
| 3 | Forward transfer   | Time-ordered: trained in 2011, scored on 2012 | Naive lag-1, lag-7 and rolling seven-day rules     | Does the approach still work on a period it has not seen? | Table 12       |

Comparisons 1 and 2 address the brief. Comparison 3 is an additional check that limits how the result may be used.

## 1.3 Evaluation criteria

MAE is primary because it reports the average miss in rentals per day. RMSE gives more weight to large errors, while R-squared shows explained variation relative to predicting the sample mean.

The training-mean baseline tests whether predictors add value; the primary model must improve MAE by 40%. Rolling seven-day demand is the temporal reference because it adapts to recent demand; the frozen ML model must beat it by 5% on identical 2012 dates.

Model, feature and hyperparameter choices use training-only cross-validation. Lowest MAE leads; configurations within 5% are separated by RMSE and then simplicity. The holdout is used only after selection.

These study-defined criteria came from earlier exploratory versions and are not stakeholder-validated service levels.

## 1.4 Assumptions and limitations

- Observed weather omits forecast error, making reported day-ahead accuracy optimistic.
- Two years provide one year-on-year transition, weak evidence for persistent growth.
- Weather category 4 is absent; an input guard must restrict use to supported categories 1-3.
- System totals cannot answer station questions. The Institute for Transportation and Development Policy (ITDP) publishes bikeshare planning guidance; the National Association of City Transportation Officials (NACTO) publishes network and station-siting guidance. They explain why future station planning also needs identifiers, capacity, trip flows, transit connections and local demand (Institute for Transportation & Development Policy [ITDP], 2018; National Association of City Transportation Officials [NACTO], 2016).

## 2. Data Understanding

### 2.1 Source, loading and granularity

The UCI Bike Sharing dataset supplies `day.csv` and `hour.csv` for Capital Bikeshare in 2011-2012 (Fanaee-T & Gama, 2014; University of California, Irvine, n.d.). The daily file is the modelling source. The hourly file is loaded separately for intraday EDA and cross-file validation.

```

In [1]: import os
import warnings
from pathlib import Path

os.environ["PYTHONWARNINGS"] = "ignore"
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.base import clone
from sklearn.dummy import DummyRegressor
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import (
    train_test_split, GridSearchCV, KFold, TimeSeriesSplit, cross_validate
)
from sklearn.inspection import permutation_importance
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

sns.set_theme(style="whitegrid")
RANDOM_STATE = 42

NB_DIR = Path.cwd()
BASE_DIR = NB_DIR.parent if NB_DIR.name == "notebook" else NB_DIR
DATA_DIR = BASE_DIR / "dataset"
OUTPUT_DIR = BASE_DIR / "outputs"
FIG_DIR = OUTPUT_DIR / "figures"
FIG_DIR.mkdir(parents=True, exist_ok=True)

day_df = pd.read_csv(DATA_DIR / "day.csv", parse_dates=["dteday"]).sort_values("
hour_df = pd.read_csv(DATA_DIR / "hour.csv", parse_dates=["dteday"]).sort_values("
print("day.csv :", day_df.shape, day_df.dteday.min().date(), "to", day_df.dteday
print("hour.csv:", hour_df.shape, hour_df.dteday.min().date(), "to", hour_df.dte

day.csv : (731, 16) 2011-01-01 to 2012-12-31
hour.csv: (17379, 17) 2011-01-01 to 2012-12-31

```

```
In [2]: hourly_daily = hour_df.groupby("dteday", as_index=False)["cnt"].sum().rename(col
cross_file = day_df[["dteday", "cnt"]].merge(hourly_daily, on="dteday", how="out
cross_file["difference"] = cross_file["hourly_sum"] - cross_file["cnt"]
assert len(cross_file) == 731
assert cross_file["difference"].eq(0).all(), "Hourly cnt does not sum to daily c
hours_per_date = hour_df.groupby("dteday").size()
incomplete_dates = int(hours_per_date.lt(24).sum())
omitted_hour_rows = int((24 - hours_per_date).sum())
assert incomplete_dates == 76 and omitted_hour_rows == 165

granularity = pd.DataFrame([
    {"file": "day.csv", "rows": len(day_df), "observation": "one daily predictio
    "role": "selected for modelling"},
    {"file": "hour.csv", "rows": len(hour_df), "observation": "correlated hourly
    "role": "intraday understanding and cross-file validation"},
])
print("Table 2 - Granularity and role")
display(granularity)
print(f"Cross-file validation passed: hourly cnt sums exactly to daily cnt on al
print(f"Hourly panel audit: {incomplete_dates} dates have fewer than 24 rows; "
      f"{omitted_hour_rows} zero-demand hour rows are omitted in total.")
```

Table 2 – Granularity and role

|   | file     | rows  | observation                    | role                                             |
|---|----------|-------|--------------------------------|--------------------------------------------------|
| 0 | day.csv  | 731   | one daily prediction unit      | selected for modelling                           |
| 1 | hour.csv | 17379 | correlated hourly observations | intraday understanding and cross-file validation |

Cross-file validation passed: hourly cnt sums exactly to daily cnt on all 731 dates.

Hourly panel audit: 76 dates have fewer than 24 rows; 165 zero-demand hour rows are omitted in total.


`hour.csv` is unbalanced: 76 of 731 dates have fewer than 24 observations, with 165 zero-demand rows omitted. An hourly model would change granularity, weight dates unequally and omit those periods. Selecting `day.csv` gives every target day one complete observation **to avoid hourly bias**. The hourly counts still reconstruct all daily totals and remain useful for intraday interpretation.

```
In [3]: hour_profile = (hour_df.groupby(["workingday", "hr"], as_index=False)["cnt"].mean)
peaks = (hour_profile.loc[hour_profile.groupby("workingday")["cnt"].idxmax()]
        .assign(day_type=lambda x: x["workingday"].map({0: "Non-working day", 1: "Working day"}))
        .rename(columns={"hr": "peak_hour", "cnt": "mean_cnt"})
peaks["mean_cnt"] = peaks["mean_cnt"].round(1)
print("Table 3 - Peak hourly demand")
display(peaks)
assert int(peaks.loc[peaks.day_type == "Working day", "peak_hour"].iloc[0]) == 17
assert np.isclose(peaks.loc[peaks.day_type == "Working day", "mean_cnt"].iloc[0], 525.3)
assert int(peaks.loc[peaks.day_type == "Non-working day", "peak_hour"].iloc[0]) == 13
assert np.isclose(peaks.loc[peaks.day_type == "Non-working day", "mean_cnt"].iloc[0], 372.7)

fig, ax = plt.subplots(figsize=(10.5, 4.6))
for flag, label, color in [(1, "Working day", "#20639B"), (0, "Non-working day", "#D62728")]:
    p = hour_profile[hour_profile.workingday == flag]
    ax.plot(p.hr, p.cnt, marker="o", ms=3, lw=2, label=label, color=color)
ax.set(title="Figure 1 - Mean hourly rentals by day type", xlabel="hour of day",
       xticks=range(0, 24, 2)); ax.legend()
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_hourly_demand.png", dpi=160); plt.
```

Table 3 - Peak hourly demand

|    | day_type        | peak_hour | mean_cnt |
|----|-----------------|-----------|----------|
| 13 | Non-working day | 13        | 372.7    |
| 41 | Working day     | 17        | 525.3    |



Working-day demand peaks at 17:00 (525.3 rentals), while non-working-day demand peaks at 13:00 (372.7). The profiles are consistent with commute-oriented and broader daytime use, respectively; aggregate counts do not prove trip purpose.

```
In [4]: span_days = (day_df.dteday.max() - day_df.dteday.min()).days + 1
norm_cols = ["temp", "atemp", "hum", "windspeed"]
quality = pd.DataFrame([
    ("Daily rows equal calendar span", len(day_df) == span_days, f"{len(day_df)}"),
    ("No date gaps", day_df.dteday.diff().dropna().dt.days.max() == 1, "daily co"),
    ("No exact duplicates", day_df.duplicated().sum() == 0, f"{day_df.duplicated"),
    ("No source missing values", day_df.isna().sum().sum() == 0, f"{day_df.isna(
    ("casual + registered = cnt", (day_df.casual + day_df.registered == day_df.c),
    ("Normalised weather in [0,1]", all(day_df[c].between(0, 1).all() for c in n
    ("Humidity above zero", (day_df.hum > 0).all(), f"{int((day_df.hum == 0).sum
    ("Daily weather categories supported", set(day_df.weathersit.unique()) == {1
]), columns=["check", "pass", "evidence"])
print("Table 4 - Data quality audit")
display(quality)

zero_hum_date = day_df.loc[day_df.hum.eq(0), "dteday"].iloc[0]
zero_hour_rows = hour_df[hour_df.dteday.eq(zero_hum_date)]
print("Zero-humidity date:", zero_hum_date.date(), "| hourly rows:", len(zero_ho
      "| hourly humidity values:", sorted(zero_hour_rows.hum.unique()))

# Keep a raw frame for sensitivity analysis, then mark the implausible value mis
raw_day_df = day_df.copy()
day_df.loc[day_df.hum.eq(0), "hum"] = np.nan
assert day_df.hum.isna().sum() == 1
```

Table 4 - Data quality audit

|   | check                              | pass  | evidence            |
|---|------------------------------------|-------|---------------------|
| 0 | Daily rows equal calendar span     | True  | 731 rows / 731 days |
| 1 | No date gaps                       | True  | daily continuity    |
| 2 | No exact duplicates                | True  | 0 duplicates        |
| 3 | No source missing values           | True  | 0 missing           |
| 4 | casual + registered = cnt          | True  | all rows            |
| 5 | Normalised weather in [0,1]        | True  | all four columns    |
| 6 | Humidity above zero                | False | 1 zero reading      |
| 7 | Daily weather categories supported | True  | category 4 absent   |

Zero-humidity date: 2011-03-10 | hourly rows: 22 | hourly humidity values: [np.float64(0.0)]

The daily series is complete and `casual + registered = cnt` holds throughout. One date reports zero humidity across 22 hourly rows. It is treated as a recording fault, changed to missing and median-imputed inside each training fold; a sensitivity check tests the decision. Plausible tail values are retained.

Weather category 4 is absent, so the supported domain is 1-3 and still requires an input guard.

## 2.2 Variables and exploratory analysis


| Variables                                                                                                                            | Role                                                                          |
|--------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| <code>cnt</code>                                                                                                                     | Continuous target: system-wide rentals per day                                |
| <code>season</code> , <code>yr</code> , <code>mnth</code> , <code>holiday</code> ,<br><code>weekday</code> , <code>workingday</code> | Calendar predictors                                                           |
| <code>weathersit</code> , <code>temp</code> , <code>atemp</code> , <code>hum</code> ,<br><code>windspeed</code>                      | Weather predictors                                                            |
| <code>dteday</code>                                                                                                                  | Ordering and elapsed-time construction; not passed directly to a model        |
| <code>casual</code> , <code>registered</code>                                                                                        | Descriptive plots only; excluded because they sum exactly to <code>cnt</code> |

`casual` and `registered` describe user groups, not independent predictors. Registered users may include commuters, but membership status does not prove trip purpose.

```
In [5]: season_lbl = {1: "winter", 2: "spring", 3: "summer", 4: "autumn"}
weather_lbl = {1: "clear", 2: "mist", 3: "light rain/snow", 4: "heavy rain/snow"}
day_df["season_name"] = day_df.season.map(season_lbl)
day_df["weather_name"] = day_df.weathersit.map(weather_lbl)
day_df["year_name"] = day_df.yr.map({0: "2011", 1: "2012"})

fig, ax = plt.subplots(2, 2, figsize=(13, 9))
sns.histplot(day_df.cnt, bins=30, kde=True, ax=ax[0, 0], color="#20639B")
ax[0, 0].set(title="Daily-rental distribution", xlabel="daily cnt")
sns.lineplot(data=day_df, x="dteday", y="cnt", hue="year_name", ax=ax[0, 1])
ax[0, 1].set(title="Daily demand over time", xlabel="date", ylabel="cnt")
sns.barplot(data=day_df, x="season_name", y="cnt", order=["winter", "spring", "summer", "autumn"], ax=ax[1, 0])
ax[1, 0].set(title="Mean demand by season", xlabel="")
sns.boxplot(data=day_df, x="weather_name", y="cnt", order=["clear", "mist", "light rain/snow", "heavy rain/snow"], ax=ax[1, 1])
ax[1, 1].set(title="Demand by weather category", xlabel="")
fig.suptitle("Figure 2 - Daily demand, season and weather", y=1.01)
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_daily_eda.png", dpi=150); plt.show
```

Figure 2 - Daily demand, season and weather



Demand rises across the two years, peaks in warmer seasons and falls in adverse weather.

**A documentation correction.** The dataset description states `season (1:summer, 2:summer, 3:fall, 4:winter)`. The dates disagree. Code 1 covers December to March at a mean of 12.2 degrees Celsius and the lowest mean demand of 2,604; code 3 covers June to September at 29.0 degrees and the highest mean demand of 5,644. The codes are therefore offset by one quarter from the published labels, and this notebook labels them `1 = winter, 2 = spring, 3 = summer, 4 = autumn` on the evidence of the calendar and temperature rather than the documentation. Nothing in the modelling changes, because `season` enters every pipeline as a categorical code and never as a name; the correction affects the readability of Figures 2 and 3 and of Table 13, and the wording of this section.

```
In [6]: comp = day_df.groupby("season_name")[["casual", "registered"]].mean().reindex(["
fig, ax = plt.subplots(figsize=(8.5, 4.4))
comp.plot(kind="bar", stacked=True, ax=ax, color=["#ED553B", "#20639B"])
ax.set(title="Figure 3 - Mean casual and registered rentals by season",
        xlabel="", ylabel="mean daily rentals")
ax.tick_params(axis="x", rotation=0)
ax.legend(title="user group")
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_composition.png", dpi=160); plt.sh

# Each weather variable is plotted in its original unit, not the shared normalis
# day.csv divides temp by 41 C, atemp by 50 C, hum by 100 % and windspeed by 67
# single normalised axis would give one x position four different physical meani
WEATHER_UNITS = [
    ("temp", 41, "temperature (degrees Celsius)", "#20639B"),
    ("atemp", 50, "feels-like temperature (degrees Celsius)", "#3CAEA3"),
    ("hum", 100, "relative humidity (%)", "#F6D55C"),
    ("windspeed", 67, "wind speed (km/h)", "#ED553B"),
]
fig, ax = plt.subplots(1, 4, figsize=(16, 4.2), sharey=True)
for axis, (column, scale, label, color) in zip(ax, WEATHER_UNITS):
    sns.regplot(x=day_df[column] * scale, y=day_df.cnt, ax=axis, order=2, ci=Non
                scatter_kws={"s": 10, "alpha": .3, "color": color},
                line_kws={"color": "black", "lw": 1.4})
    axis.set(xlabel=label, ylabel="")
ax[0].set(ylabel="daily cnt")
fig.suptitle("Figure 4 - Weather measures and daily demand, each in its original
              "(quadratic guide line, not a fitted model)", y=1.04)
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_weather_units.png", dpi=160, bbox_
```





Figure 4 - Weather measures and daily demand, each in its original unit (quadratic guide line, not a fitted model)



Registered users provide most rentals in every season, while the casual share expands in the warmer ones. Both columns stay descriptive because they sum exactly to the target.


Plotting each weather variable on its own axis in real units matters for reading the strength of each relationship. Demand climbs steeply with temperature to roughly 30 degrees Celsius and then stops climbing, which is the curvature the pairplot also shows and the reason a squared temperature term is built in Section 3.3. Feels-like temperature repeats the same shape, as its correlation of +0.631 against +0.627 would suggest.

The other two behave differently from each other, which the shared axis in the earlier version of this figure concealed. Wind speed declines steadily across its range. Humidity does not: the guide line peaks near 60% and falls away at both ends, so the weak -0.115 correlation in Table 5 is summarising a non-monotonic pattern rather than a gentle slope, and the extremes are thinly populated. Neither variable separates high-demand from low-demand days on its own; the scatter at any given humidity or wind speed spans most of the observed range of `cnt`.

A single normalised axis would have hidden all of this by placing 20.5 degrees Celsius, 25 degrees Celsius, 50% humidity and 33.5 km/h at the same x position.

```
In [7]: weekday_lbl = {0: "Sun", 1: "Mon", 2: "Tue", 3: "Wed", 4: "Thu", 5: "Fri", 6: "Sat"}
        day_df["weekday_name"] = day_df.weekday.map(weekday_lbl)

        fig, ax = plt.subplots(1, 2, figsize=(13, 4.5))
        sns.lineplot(data=day_df, x="mnth", y="cnt", hue="year_name", marker="o", ax=ax[0])
        ax[0].set(title="Figure 5 - Mean rentals by month and year", xlabel="month", ylabel="mean cnt")
        ax[0].set_xticks(range(1, 13))
        sns.barplot(data=day_df, x="weekday_name", y="cnt",
                    order=["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"], ax=ax[1])
        ax[1].set(title="Figure 6 - Mean rentals by weekday", xlabel="weekday", ylabel="mean cnt")
        plt.tight_layout(); plt.savefig(FIG_DIR / "v5_month_weekday.png", dpi=160); plt.
```



Monthly demand follows the seasonal wave and every 2012 month exceeds its 2011 counterpart. Weekday means vary much less, showing that aggregation blends working-day and leisure patterns visible in the hourly profiles.

```
In [8]: corr_cols = ["season", "yr", "mnth", "holiday", "weekday", "workingday",
                    "weathersit", "temp", "atemp", "hum", "windspeed", "cnt"]
corr_matrix = day_df[corr_cols].corr()
cnt_correlations = corr_matrix["cnt"].drop("cnt").sort_values(ascending=False)
assert cnt_correlations.index[0] == "atemp"
assert cnt_correlations.loc["weathersit"] < 0 and cnt_correlations.loc["windspeed"] > 0

plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", center=0,
            square=True, cbar_kws={"shrink": .8}, annot_kws={"size": 7})
plt.title("Figure 7 - Correlation matrix for daily variables")
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_correlation_heatmap.png", dpi=160)
print("Table 5 - Correlations with cnt")
display(cnt_correlations.round(3).to_frame("correlation_with_cnt"))
```





Table 5 - Correlations with cnt

|                   | correlation_with_cnt |
|-------------------|----------------------|
| <b>atemp</b>      | 0.631                |
| <b>temp</b>       | 0.627                |
| <b>yr</b>         | 0.567                |
| <b>season</b>     | 0.406                |
| <b>mnth</b>       | 0.280                |
| <b>weekday</b>    | 0.067                |
| <b>workingday</b> | 0.061                |
| <b>holiday</b>    | -0.068               |
| <b>hum</b>        | -0.115               |
| <b>windspeed</b>  | -0.235               |
| <b>weathersit</b> | -0.297               |

`atemp` has the highest correlation with `cnt` (+0.631), followed by `temp` (+0.627) and `yr` (+0.567). Adverse weather ( `weathersit` , -0.297), windspeed (-0.235) and corrected humidity (-0.115) are negative. The near-collinearity of `temp` and `atemp` cautions against interpreting linear coefficients independently; categorical codes also limit causal interpretation.

```
In [9]: # Stage 2.6: use seaborn's pairplot function and include its output.
pairplot_data = day_df[["temp", "atemp", "hum", "windspeed", "cnt"]].dropna()
pair_grid = sns.pairplot(pairplot_data, corner=True, diag_kind="kde",
                        plot_kws={"alpha": .3, "s": 12})
pair_grid.fig.suptitle("Figure 8 - Seaborn pairplot of weather variables and dai
pair_grid.savefig(FIG_DIR / "v5_pairplot.png", dpi=140, bbox_inches="tight")
plt.show()
```

Figure 8 - Seaborn pairplot of weather variables and daily demand



The pairplot confirms positive, nonlinear relationships between demand and both temperature measures. Humidity and windspeed show weaker negative patterns and broad scatter. `temp` and `atemp` move almost identically, consistent with their high correlation, while no single weather variable fully explains demand.

## 3. Data Preparation

Section 2 reported what the data contains. This section records the decisions taken because of it, and the reasoning behind each one.

### 3.1 Selecting the data

**Daily rather than hourly records.** Both files describe the same system, so the choice of granularity is a modelling decision and not a formatting one. Section 2.1 showed that `hour.csv` is an unbalanced panel: 76 of the 731 dates carry fewer than 24 rows, and the 165 absent rows are zero-demand hours that were never written. Fitting on hourly records would therefore weight dates unequally, giving a complete summer date 24 observations and an incomplete winter date as few as 22, and it would never show the model the quiet hours that are missing. Selecting `day.csv` gives every target day exactly one complete observation and **avoids that hourly bias**. The hourly file is retained for intraday interpretation and for the cross-file check that reconstructs all 731 daily totals exactly.

**Columns excluded on principle.** `casual` and `registered` are dropped from every predictor list because they sum to `cnt` by construction; including either would let the model read part of its own answer. `dteday` is not passed to any estimator either, and is used only for ordering and for building elapsed time. `instant` is a row counter carrying no information about demand.

**Two frames, so the temporal question does not shrink the primary sample.**

`primary_df` keeps all 731 days and contains no past-demand inputs. `temporal_df` begins after a causal seven-day warm-up and contains 724 days, because a seven-day rolling mean is undefined for the first week of 2011. Building one shared frame would have cost seven days from the primary experiment for no benefit to it.

### 3.2 Cleaning, and estimating the missing value

The daily file arrives with no missing cells, no duplicate rows and no gaps in the date sequence, so cleaning reduces to a single judgement call.

**One humidity reading is not a low value, it is a fault.** Section 2.1 recorded `hum = 0` on one date. Relative humidity of exactly zero does not occur in Washington DC, and cross-checking `hour.csv` confirms the diagnosis: all 22 hourly rows for that date also read zero, which is the signature of a failed sensor rather than of dry weather.

**The value is estimated, not deleted.** Deleting the row would have discarded a complete demand observation to fix one input, so the reading is marked missing and then estimated by the median of the training data. The estimation happens inside `SimpleImputer(strategy="median")` within every pipeline, so the median is refitted on each training fold and never sees the validation fold or the holdout. A single `fillna` computed over the whole frame would have been simpler and would have

leaked a statistic of the test data into a training input. The interaction term `atemp_hum` inherits the correction automatically, because it is rebuilt from the corrected column.

**The decision is checked rather than asserted.** Table 9 refits the frozen primary model on the uncorrected data and reports both results. If the correction had been load-bearing, that table would say so.

**No outlier removal.** Extreme but physically plausible values are retained. Low-demand days are real events, mostly adverse weather, and removing them would flatter the error statistics while making the model worse at exactly the days a planner cares about.

### 3.3 Engineering, encoding and the split

**What the correlations imply for preparation.** `atemp` correlates with `cnt` at +0.631 and `temp` at +0.627, and the two move almost identically with each other. Both are retained, because the tree families are untroubled by collinear inputs and dropping one would discard information the ensembles use. The consequence is recorded instead: individual linear coefficients on `temp` and `atemp` are not interpretable in isolation, which is one reason Section 5.2 explains the model through permutation importance and TreeSHAP rather than through coefficients. `weathersit` (-0.297), `windspeed` (-0.235) and corrected `hum` (-0.115) are the negative correlates, and all three are kept as predictors.

**Interaction and curvature terms.** `atemp * hum` and `temp` squared are constructed so that the additive families can express what the scatter plots show: warm and humid days behave differently from warm and dry ones, and demand stops rising once temperature passes roughly 30 degrees Celsius. Both are built only from target-day exogenous inputs, so neither smuggles in past demand.

**Cyclical encoding, in the temporal experiment only.** Month, weekday and season are also encoded as sine and cosine pairs for the time-ordered sets. The reason is specific to that split: `OneHotEncoder(drop="first", handle_unknown="ignore")` encodes an unseen category as all zeros, which is byte-identical to the dropped reference category. Under a 2011-to-2012 split, a month the model never met would silently be treated as January. A sine and cosine pair has no unseen values and keeps December adjacent to January.

**`yr` is excluded from the temporal predictors.** It is constant at 0 across the whole 2011 training year, so the model can learn nothing from it, and it then arrives as a constant 1 in 2012, a value never observed during fitting. It stays in the primary feature sets, where both values appear on both sides of the split.

**Autoregressive inputs are built causally.** `lag_1_cnt`, `lag_7_cnt` and `roll_7_cnt` use only demand observed strictly before the target day; the rolling mean shifts first and then rolls, never the reverse. The code cell below asserts this for every row rather than trusting the expression, because shift-after-roll is the easiest way to leak a target into its own predictor.

**The split, declared here and executed in Section 4.** The primary experiment uses a single random 75/25 partition created once at `random_state=42`, giving 548 training and 183 holdout days. Selection of family, feature set and hyperparameters happens entirely inside the 548 training rows by cross-validation; the 183 holdout rows are scored once, after everything is frozen. The temporal experiment splits by calendar instead, fitting on 2011 and scoring on 2012.

```
In [10]: # Shared calendar/weather engineering. Interactions use only target-day exogenous
day_df["atemp_hum"] = day_df["atemp"] * day_df["hum"]
day_df["temp_sq"] = day_df["temp"] ** 2
day_df["days_since_start"] = (day_df.dteday - day_df.dteday.min()).dt.days
day_df["mnth_sin"], day_df["mnth_cos"] = np.sin(2*np.pi*day_df.mnth/12), np.cos(2*np.pi*day_df.mnth/12)
day_df["weekday_sin"], day_df["weekday_cos"] = np.sin(2*np.pi*day_df.weekday/7), np.cos(2*np.pi*day_df.weekday/7)
day_df["season_sin"], day_df["season_cos"] = np.sin(2*np.pi*day_df.season/4), np.cos(2*np.pi*day_df.season/4)

primary_df = day_df.copy().reset_index(drop=True)

# Causal temporal features: every window ends before the target day.
day_df["lag_1_cnt"] = day_df.cnt.shift(1)
day_df["lag_7_cnt"] = day_df.cnt.shift(7)
day_df["roll_7_cnt"] = day_df.cnt.shift(1).rolling(7).mean()
assert day_df.lag_1_cnt.equals(day_df.cnt.shift(1))
assert day_df.lag_7_cnt.equals(day_df.cnt.shift(7))
for i in range(7, len(day_df)):
    assert np.isclose(day_df.loc[i, "roll_7_cnt"], day_df.loc[i-7:i-1, "cnt"].mean())
temporal_df = day_df.dropna(subset=["roll_7_cnt"]).reset_index(drop=True)

assert len(primary_df) == 731
assert len(temporal_df) == 724
print("primary_df rows:", len(primary_df), "| temporal_df rows:", len(temporal_df))
```

primary\_df rows: 731 | temporal\_df rows: 724

```

In [11]: PRIMARY_NOMINAL = ["season", "mnth", "weekday", "weathersit"]
PRIMARY_BINARY = ["yr", "holiday", "workingday"]
WEATHER = ["temp", "atemp", "hum", "windspeed"]
CYCLICAL = ["mnth_sin", "mnth_cos", "weekday_sin", "weekday_cos", "season_sin",
INTERACTIONS = ["atemp_hum", "temp_sq"]
TREND = ["days_since_start"]
AUTOREGRESSIVE = ["lag_1_cnt", "lag_7_cnt", "roll_7_cnt"]

PRIMARY_FEATURE_SETS = {
    "core": PRIMARY_NOMINAL + PRIMARY_BINARY + WEATHER,
    "core + interactions": PRIMARY_NOMINAL + PRIMARY_BINARY + WEATHER + INTERACTIONS,
    "core + interactions + trend": PRIMARY_NOMINAL + PRIMARY_BINARY + WEATHER + TREND,
}
TEMPORAL_FEATURE_SETS = {
    "core": ["weathersit", "holiday", "workingday"] + WEATHER + CYCLICAL,
    "core + interactions": ["weathersit", "holiday", "workingday"] + WEATHER + INTERACTIONS,
    "core + interactions + autoregressive": ["weathersit", "holiday", "workingday"] + WEATHER + AUTOREGRESSIVE,
    "core + interactions + autoregressive + trend": ["weathersit", "holiday", "workingday"] + WEATHER + AUTOREGRESSIVE + TREND,
}

for feature_list in list(PRIMARY_FEATURE_SETS.values()) + list(TEMPORAL_FEATURE_SETS.values()):
    assert "casual" not in feature_list and "registered" not in feature_list
assert all(c not in cols for cols in PRIMARY_FEATURE_SETS.values() for c in AUTOREGRESSIVE)
assert all("yr" not in cols for cols in TEMPORAL_FEATURE_SETS.values())

def make_preprocessor(columns, temporal=False):
    nominal = ["weathersit"] if temporal else [c for c in PRIMARY_NOMINAL if c in columns]
    binary_pool = ["holiday", "workingday"] + ([ "yr" ] if temporal else [])
    binary = [c for c in binary_pool if c in columns]
    continuous = [c for c in columns if c not in nominal + binary]
    return ColumnTransformer([
        ("nominal", OneHotEncoder(drop="first", handle_unknown="ignore"), nominal),
        ("continuous", Pipeline([("impute", SimpleImputer(strategy="median")),
                                ("scale", StandardScaler())]), continuous),
        ("binary", "passthrough", binary),
    ])

def pipeline_for(model, columns, temporal=False):
    return Pipeline([("prep", make_preprocessor(columns, temporal=temporal)), (model_name, model)])

def metrics(y_true, predictions):
    mse = mean_squared_error(y_true, predictions)
    return {"MAE": mean_absolute_error(y_true, predictions), "MSE": mse,
            "RMSE": float(np.sqrt(mse)), "R2": r2_score(y_true, predictions)}

GRIDS = {
    "Linear Regression": (LinearRegression(), {}),
    "K-Nearest Neighbors": (KNeighborsRegressor(), {
        "model__n_neighbors": [3, 5, 7, 9, 11, 15], "model__weights": ["uniform", "distance"]
    }),
    "Random Forest": (RandomForestRegressor(random_state=RANDOM_STATE, n_jobs=-1), {
        "model__n_estimators": [200, 400], "model__max_depth": [None, 8, 16],
        "model__min_samples_leaf": [1, 2, 4]}),
    "Gradient Boosting": (GradientBoostingRegressor(random_state=RANDOM_STATE), {
        "model__n_estimators": [200, 400], "model__max_depth": [2, 3],
        "model__learning_rate": [0.05, 0.1]}),
}

MODEL_SIMPLICITY = {"Linear Regression": 1, "K-Nearest Neighbors": 2,
                    "Random Forest": 3, "Gradient Boosting": 4}
FEATURE_SIMPLICITY = {name: i for i, name in enumerate(
    ["core", "core + interactions", "core + interactions + trend",
    "core + interactions + autoregressive", "core + interactions + autoregressive + trend"])}

def compact_params(params):
    if not params:

```

```

        return "default"
    labels = {"model__learning_rate": "lr", "model__max_depth": "depth",
              "model__n_estimators": "estimators", "model__min_samples_leaf": "l",
              "model__n_neighbors": "neighbors", "model__weights": "weights"}
    return "; ".join(f"{labels.get(key, key)}={value}" for key, value in params.

def choose_configuration(rows):
    frame = pd.DataFrame(rows).copy()
    best_mae = frame["cv_mean_MAE"].min()
    shortlist = frame[frame["cv_mean_MAE"] <= 1.05 * best_mae].copy()
    shortlist["simplicity"] = shortlist["model"].map(MODEL_SIMPLICITY) + shortli
    return shortlist.sort_values(["cv_mean_RMSE", "simplicity", "cv_mean_MAE"]).

```

Preparation leaves 731 primary rows and 724 temporal rows, with no observation deleted and one input value estimated. The assertions above are part of the deliverable rather than scaffolding: they prove that the rolling window closes before the target day, that no predictor list contains `casual` or `registered`, that past-demand inputs never reach the primary experiment, and that `yr` never reaches the temporal one. Each is a failure this design could plausibly have suffered, so each is checked by the notebook instead of being claimed by the report.

Every transformation that learns anything from the data, meaning median imputation and standard scaling, is defined inside the pipeline and not applied to the frame beforehand. That placement is what allows Section 4 to refit the whole chain on each cross-validation fold without a training statistic escaping into evaluation data.

## 4. Modelling

Four regression families are compared with lightweight grids in scikit-learn (Pedregosa et al., 2011). Linear Regression tests a global additive relationship; K-Nearest Neighbors tests local similarity; Random Forest averages decorrelated trees; Gradient Boosting sequentially corrects residuals (Friedman, 2001). Scaling and fold-fitted median imputation live inside every pipeline.

Both experiments run over the same 731 days and differ in how those days are split and which inputs are permitted. Each produces one learned estimate and one reference estimate, giving the four columns compared in Table 13.

```
In [12]: from IPython.display import Image, display as show_image

# Figure 9 is authored in Mermaid and rendered to PNG so that it survives the no
# HTML and PDF exports identically. The source is kept in the repository alongsi
show_image(Image(filename=str(FIG_DIR / "v5_pipeline.png"), width=760))
print("Figure 9 - How the two experiments produce the four estimates compared in
```




Figure 9 – How the two experiments produce the four estimates compared in Table 1  
3

## 4.1 Primary model comparison: random 75/25 holdout

The split is created once with `random_state=42`. Only the 75% training partition participates in family, feature-set and hyperparameter selection. Each of the 12 family/feature-set combinations uses shuffled five-fold `KFold` and MAE scoring. The 25% holdout is accessed only after the global configuration and one configuration per family have been frozen.

```

In [13]: all_indices = primary_df.index.to_numpy()
primary_train_idx, primary_holdout_idx = train_test_split(
    all_indices, test_size=0.25, random_state=RANDOM_STATE
)
assert set(primary_train_idx).isdisjoint(primary_holdout_idx)
assert len(primary_train_idx) == 548 and len(primary_holdout_idx) == 183

primary_train = primary_df.loc[primary_train_idx].copy()
primary_holdout = primary_df.loc[primary_holdout_idx].copy()
primary_cv = KFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)
primary_rows, primary_estimators = [], {}

for feature_name, columns in PRIMARY_FEATURE_SETS.items():
    for model_name, (estimator, grid) in GRIDS.items():
        search = GridSearchCV(
            pipeline_for(clone(estimator), columns, temporal=False), grid,
            scoring="neg_mean_absolute_error", cv=primary_cv, n_jobs=1, refit=True
        )
        search.fit(primary_train[columns], primary_train.cnt)
        rmse_cv = cross_validate(
            search.best_estimator_, primary_train[columns], primary_train.cnt,
            cv=primary_cv, scoring="neg_root_mean_squared_error", n_jobs=1
        )["test_score"]
        row = {
            "model": model_name, "feature_set": feature_name,
            "cv_mean_MAE": -search.best_score_,
            "cv_std_MAE": search.cv_results_["std_test_score"][search.best_index],
            "cv_mean_RMSE": -rmse_cv.mean(), "cv_std_RMSE": rmse_cv.std(),
            "best_params": compact_params(search.best_params_),
        }
        primary_rows.append(row)
        primary_estimators[(model_name, feature_name)] = search.best_estimator_

primary_cv_table = (pd.DataFrame(primary_rows)
                    .sort_values(["cv_mean_MAE", "cv_mean_RMSE"])
                    .reset_index(drop=True))
primary_cv_table.to_csv(OUTPUT_DIR / "primary_cv_selection_v5.csv", index=False)

primary_choice = choose_configuration(primary_rows)
primary_family_choices = {
    family: choose_configuration([r for r in primary_rows if r["model"] == family])
    for family in GRIDS
}
primary_winner_name = primary_choice["model"]
primary_winner_features = primary_choice["feature_set"]
primary_winner_columns = PRIMARY_FEATURE_SETS[primary_winner_features]
primary_winner = primary_estimators[(primary_winner_name, primary_winner_features)]

# Leakage audit: choices contain CV statistics from training indices only; the h
# been passed to fit, GridSearchCV, cross_validate or choose_configuration.
selection_indices = set(primary_train_idx)
assert selection_indices.isdisjoint(primary_holdout_idx)
assert len(primary_df) == 731
print("Table 6 – Primary CV selection (training partition only)")
display(primary_cv_table.round({"cv_mean_MAE": 1, "cv_std_MAE": 1,
                                "cv_mean_RMSE": 1, "cv_std_RMSE": 1}))
print("Frozen primary configuration:", primary_winner_name, "|", primary_winner_
      "|", primary_choice["best_params"])

```

Table 6 – Primary CV selection (training partition only)

|    | model               | feature_set                 | cv_mean_MAE | cv_std_MAE | cv_mean_RMSE | cv_std_RMSE | best_params                                         |
|----|---------------------|-----------------------------|-------------|------------|--------------|-------------|-----------------------------------------------------|
| 0  | Gradient Boosting   | core + interactions + trend | 477.0       | 34.1       | 657.6        | 45.4        | lr=0.1; depth=2; estimators=400                     |
| 1  | Random Forest       | core + interactions + trend | 498.0       | 18.1       | 723.5        | 38.5        | depth=16; lr=0.05; estimators=400                   |
| 2  | Gradient Boosting   | core                        | 513.8       | 56.4       | 690.8        | 53.2        | lr=0.05; depth=2; estimators=400                    |
| 3  | Gradient Boosting   | core + interactions         | 519.5       | 47.4       | 697.1        | 46.4        | lr=0.05; depth=2; estimators=400                    |
| 4  | Random Forest       | core                        | 533.6       | 26.5       | 746.8        | 25.7        | depth=16; lr=0.05; estimators=400                   |
| 5  | Random Forest       | core + interactions         | 539.1       | 24.3       | 752.6        | 23.2        | depth=16; lr=0.05; estimators=400                   |
| 6  | Linear Regression   | core + interactions + trend | 559.3       | 19.8       | 752.4        | 36.9        | coef=[0.00000000e+00 1.00000000e+00 0.00000000e+00] |
| 7  | Linear Regression   | core + interactions         | 560.4       | 19.3       | 752.6        | 38.5        | coef=[0.00000000e+00 1.00000000e+00 0.00000000e+00] |
| 8  | K-Nearest Neighbors | core + interactions + trend | 563.0       | 15.1       | 794.0        | 40.0        | neighbors=10; weights=distance                      |
| 9  | Linear Regression   | core                        | 591.0       | 14.7       | 806.8        | 41.9        | coef=[0.00000000e+00 1.00000000e+00 0.00000000e+00] |
| 10 | K-Nearest Neighbors | core + interactions         | 756.5       | 39.1       | 966.5        | 48.0        | neighbors=10; weights=distance                      |
| 11 | K-Nearest Neighbors | core                        | 760.7       | 42.7       | 957.3        | 52.9        | neighbors=10; weights=distance                      |

Frozen primary configuration: Gradient Boosting | core + interactions + trend | lr=0.1; depth=2; estimators=400

```

In [14]: # First use of the fixed 25% holdout: score one training-selected configuration
primary_holdout_rows = []
mean_model = DummyRegressor(strategy="mean").fit(
    np.zeros((len(primary_train), 1)), primary_train.cnt
)
baseline_pred = mean_model.predict(np.zeros((len(primary_holdout), 1)))
primary_holdout_rows.append({"model": "Mean baseline", "feature_set": "constant",
                             "best_params": "default", **metrics(primary_holdout

primary_family_predictions = {}
for family, choice in primary_family_choices.items():
    feature_name = choice["feature_set"]
    columns = PRIMARY_FEATURE_SETS[feature_name]
    fitted = primary_estimators[(family, feature_name)]
    pred = fitted.predict(primary_holdout[columns])
    primary_family_predictions[family] = pred
    primary_holdout_rows.append({"model": family, "feature_set": feature_name,
                                 "best_params": choice["best_params"],
                                 **metrics(primary_holdout.cnt, pred)})

primary_holdout_table = (pd.DataFrame(primary_holdout_rows)
                          .sort_values("MAE").reset_index(drop=True))
primary_holdout_table.to_csv(OUTPUT_DIR / "primary_holdout_metrics_v5.csv", index=False)
primary_metrics = primary_holdout_table.loc[
    primary_holdout_table.model.eq(primary_winner_name)
].iloc[0]
primary_baseline = primary_holdout_table.loc[
    primary_holdout_table.model.eq("Mean baseline")
].iloc[0]
primary_improvement = 1 - primary_metrics.MAE / primary_baseline.MAE
primary_mae_pct = primary_metrics.MAE / primary_holdout.cnt.mean()
primary_summary = pd.DataFrame([
    "selected_model": primary_winner_name,
    "feature_set": primary_winner_features,
    "best_params": primary_choice["best_params"],
    "holdout_MAE": primary_metrics.MAE,
    "holdout_RMSE": primary_metrics.RMSE,
    "holdout_R2": primary_metrics.R2,
    "mean_baseline_MAE": primary_baseline.MAE,
    "baseline_improvement_pct": 100 * primary_improvement,
    "MAE_pct_of_holdout_mean_demand": 100 * primary_mae_pct,
])
primary_summary.to_csv(OUTPUT_DIR / "primary_summary_v5.csv", index=False)
print("Table 7 - Frozen family configurations on the primary holdout")
display(primary_holdout_table.round({"MAE": 1, "MSE": 1, "RMSE": 1, "R2": 3}))
print("Table 8 - Final selected-model summary")
display(primary_summary.round(3))

```

Table 7 - Frozen family configurations on the primary holdout

|   | model               | feature_set                 | best_params                      | MAE    | MSE       | RMSE   | R2     |
|---|---------------------|-----------------------------|----------------------------------|--------|-----------|--------|--------|
| 0 | Gradient Boosting   | core + interactions + trend | lr=0.1; depth=2; estimators=400  | 433.9  | 397608.0  | 630.6  | 0.897  |
| 1 | Random Forest       | core + interactions + trend | depth=16; leaf=1; estimators=200 | 446.6  | 463735.1  | 681.0  | 0.880  |
| 2 | K-Nearest Neighbors | core + interactions + trend | neighbors=7; weights=distance    | 501.0  | 612073.9  | 782.4  | 0.841  |
| 3 | Linear Regression   | core + interactions + trend | default                          | 529.9  | 524664.9  | 724.3  | 0.864  |
| 4 | Mean baseline       | constant training mean      | default                          | 1663.0 | 3933577.2 | 1983.3 | -0.021 |

Table 8 - Final selected-model summary

|   | selected_model    | feature_set                 | best_params                     | holdout_MAE | holdout_RMSE | holdout_R2 | mean_b |
|---|-------------------|-----------------------------|---------------------------------|-------------|--------------|------------|--------|
| 0 | Gradient Boosting | core + interactions + trend | lr=0.1; depth=2; estimators=400 | 433.863     | 630.562      | 0.897      |        |

```
In [15]: # Sensitivity of the frozen primary winner to retaining the raw hum=0 record.
raw_primary = primary_df.copy()
raw_primary.loc[raw_primary.dteday.eq(zero_hum_date), "hum"] = 0.0
raw_primary["atemp_hum"] = raw_primary.atemp * raw_primary.hum
sens_model = clone(primary_winner).fit(
    raw_primary.loc[primary_train_idx, primary_winner_columns],
    raw_primary.loc[primary_train_idx, "cnt"]
)
sens_pred = sens_model.predict(raw_primary.loc[primary_holdout_idx, primary_winner_columns])
corrected_pred = primary_winner.predict(primary_holdout[primary_winner_columns])
humidity_sensitivity = pd.DataFrame([
    {"treatment": "hum=0 marked missing; fold-fitted median imputation",
     **metrics(primary_holdout.cnt, corrected_pred)},
    {"treatment": "raw hum=0 retained", **metrics(primary_holdout.cnt, sens_pred)}
])
humidity_sensitivity.to_csv(OUTPUT_DIR / "humidity_sensitivity_v5.csv", index=False)
print("Table 9 – Humidity correction sensitivity on the frozen primary model")
display(humidity_sensitivity.round({"MAE": 2, "MSE": 1, "RMSE": 2, "R2": 4}))
```

Table 9 – Humidity correction sensitivity on the frozen primary model

|   | treatment                                           | MAE    | MSE      | RMSE   | R2     |
|---|-----------------------------------------------------|--------|----------|--------|--------|
| 0 | hum=0 marked missing; fold-fitted median imputation | 433.86 | 397608.0 | 630.56 | 0.8968 |
| 1 | raw hum=0 retained                                  | 430.85 | 397549.6 | 630.52 | 0.8968 |

**Primary result.** Gradient Boosting with interactions and elapsed-time trend was frozen at CV MAE 477.0 (SD 34.1) and RMSE 657.6. Random Forest entered the 5% shortlist at MAE 498.0 but had higher RMSE (723.5). The selected grid point used learning rate 0.1, depth 2 and 400 estimators.

On the untouched 183-day holdout it reached MAE 433.9, RMSE 630.6 and R-squared 0.897. MAE was 73.9% below the baseline's 1,663.0 and equalled **10.1% of mean holdout demand**, meeting the 40% criterion. Random Forest reached 446.6, KNN 501.0 and Linear Regression 529.9; the winner remained frozen.

Marking zero humidity missing produced MAE 433.9 versus 430.9 when retained. The sub-1% difference leaves the conclusion unchanged; fold-fitted imputation preserves the physically plausible treatment.

## 4.2 Secondary temporal robustness check

Selection now uses `TimeSeriesSplit(n_splits=3)` inside 2011. The temporal design excludes `yr`, which is constant in the training year. It compares core cyclical calendar/weather inputs, interactions, causal recent-demand inputs and an optional elapsed-time trend. One configuration per family and the global temporal configuration are frozen before any 2012 score is produced.

```

In [16]: temporal_train = temporal_df[temporal_df.yr.eq(0)].copy().reset_index(drop=True)
temporal_test = temporal_df[temporal_df.yr.eq(1)].copy().reset_index(drop=True)
assert temporal_train.dteday.max() < pd.Timestamp("2012-01-01")
assert temporal_test.dteday.min() >= pd.Timestamp("2012-01-01")
assert all("yr" not in columns for columns in TEMPORAL_FEATURE_SETS.values())

temporal_cv = TimeSeriesSplit(n_splits=3)
fold_table = []
for fold, (tr, va) in enumerate(temporal_cv.split(temporal_train), 1):
    fold_table.append({
        "fold": fold, "train_rows": len(tr), "validation_rows": len(va),
        "train_window": f"{temporal_train.dteday.iloc[tr[0]].date()} to {temporal_train.dteday.iloc[tr[-1]].date()}",
        "validation_window": f"{temporal_train.dteday.iloc[va[0]].date()} to {temporal_train.dteday.iloc[va[-1]].date()}"
    })
print("Table 10 - TimeSeriesSplit folds within 2011")
display(pd.DataFrame(fold_table))

temporal_rows, temporal_estimators = [], {}
for feature_name, columns in TEMPORAL_FEATURE_SETS.items():
    for model_name, (estimator, grid) in GRIDS.items():
        search = GridSearchCV(
            pipeline_for(clone(estimator), columns, temporal=True), grid,
            scoring="neg_mean_absolute_error", cv=temporal_cv, n_jobs=1, refit=True
        )
        search.fit(temporal_train[columns], temporal_train.cnt)
        rmse_cv = cross_validate(
            search.best_estimator_, temporal_train[columns], temporal_train.cnt,
            cv=temporal_cv, scoring="neg_root_mean_squared_error", n_jobs=1
        )["test_score"]
        row = {"model": model_name, "feature_set": feature_name,
            "cv_mean_MAE": -search.best_score_,
            "cv_std_MAE": search.cv_results_["std_test_score"][search.best_index_],
            "cv_mean_RMSE": -rmse_cv.mean(), "cv_std_RMSE": rmse_cv.std(),
            "best_params": compact_params(search.best_params_)}
        temporal_rows.append(row)
        temporal_estimators[(model_name, feature_name)] = search.best_estimator_

temporal_cv_table = (pd.DataFrame(temporal_rows)
    .sort_values(["cv_mean_MAE", "cv_mean_RMSE"]).reset_index(drop=True))
temporal_cv_table.to_csv(OUTPUT_DIR / "temporal_cv_selection_v5.csv", index=False)
temporal_choice = choose_configuration(temporal_rows)
temporal_family_choices = {
    family: choose_configuration([r for r in temporal_rows if r["model"] == family])
    for family in GRIDS
}
temporal_winner_name = temporal_choice["model"]
temporal_winner_features = temporal_choice["feature_set"]
temporal_winner = temporal_estimators[(temporal_winner_name, temporal_winner_features)]

# Selection records and fitted estimators contain 2011 only. 2012 has not been seen
assert temporal_train.yr.eq(0).all()
assert not temporal_train.dteday.isin(temporal_test.dteday).any()
print("Table 11 - Temporal CV selection (2011 only)")
display(temporal_cv_table.round({"cv_mean_MAE": 1, "cv_std_MAE": 1,
    "cv_mean_RMSE": 1, "cv_std_RMSE": 1}))
print("Frozen temporal configuration:", temporal_winner_name, "|", temporal_winner_features,
    "|", temporal_choice["best_params"])

```

Table 10 – TimeSeriesSplit folds within 2011

|   | fold | train_rows | validation_rows | train_window             | validation_window        |
|---|------|------------|-----------------|--------------------------|--------------------------|
| 0 | 1    | 91         | 89              | 2011-01-08 to 2011-04-08 | 2011-04-09 to 2011-07-06 |
| 1 | 2    | 180        | 89              | 2011-01-08 to 2011-07-06 | 2011-07-07 to 2011-10-03 |
| 2 | 3    | 269        | 89              | 2011-01-08 to 2011-10-03 | 2011-10-04 to 2011-12-31 |

Table 11 – Temporal CV selection (2011 only)

|    | model               | feature_set                                  | cv_mean_MAE | cv_std_MAE | cv_mean_RMSE | cv_std_RMSE | best_      |
|----|---------------------|----------------------------------------------|-------------|------------|--------------|-------------|------------|
| 0  | Linear Regression   | core + interactions + autoregressive         | 602.5       | 68.5       | 763.2        | 82.4        |            |
| 1  | Linear Regression   | core + interactions + autoregressive + trend | 665.8       | 49.4       | 815.2        | 50.8        |            |
| 2  | Linear Regression   | core + interactions                          | 763.9       | 128.1      | 929.9        | 96.0        |            |
| 3  | Gradient Boosting   | core                                         | 878.0       | 356.6      | 1008.2       | 341.8       | lr=0.1; c  |
| 4  | Gradient Boosting   | core + interactions                          | 889.5       | 371.3      | 1025.5       | 351.1       | lr=0.1; c  |
| 5  | Gradient Boosting   | core + interactions + autoregressive + trend | 900.0       | 459.1      | 1086.4       | 406.9       | lr=0.05; c |
| 6  | Gradient Boosting   | core + interactions + autoregressive         | 905.0       | 459.5      | 1059.6       | 425.8       | lr=0.1; c  |
| 7  | Linear Regression   | core                                         | 917.3       | 321.5      | 1072.1       | 299.1       |            |
| 8  | Random Forest       | core + interactions + autoregressive + trend | 922.1       | 410.7      | 1094.3       | 368.3       | dept       |
| 9  | Random Forest       | core + interactions + autoregressive         | 924.4       | 401.3      | 1089.4       | 357.7       | dept       |
| 10 | Random Forest       | core                                         | 962.3       | 386.2      | 1111.3       | 355.9       | depth=8    |
| 11 | Random Forest       | core + interactions                          | 974.3       | 381.9      | 1121.1       | 352.8       | dept       |
| 12 | K-Nearest Neighbors | core                                         | 1056.4      | 349.3      | 1269.4       | 281.8       | neigh      |
| 13 | K-Nearest Neighbors | core + interactions                          | 1068.6      | 335.8      | 1269.9       | 266.2       | weights=   |
| 14 | K-Nearest Neighbors | core + interactions + autoregressive + trend | 1093.2      | 553.9      | 1314.7       | 483.4       | neigh      |
| 15 | K-Nearest Neighbors | core + interactions + autoregressive         | 1123.7      | 535.1      | 1327.1       | 476.2       | weights=   |

Frozen temporal configuration: Linear Regression | core + interactions + autoregressive | default

```
In [17]: # First 2012 scoring: frozen family candidates and causal baselines use identical
temporal_rows_2012, temporal_predictions = [], {}
for family, choice in temporal_family_choices.items():
    feature_name = choice["feature_set"]
    columns = TEMPORAL_FEATURE_SETS[feature_name]
    fitted = temporal_estimators[(family, feature_name)]
    pred = fitted.predict(temporal_test[columns])
    temporal_predictions[family] = pred
    temporal_rows_2012.append({"model": family, "feature_set": feature_name,
                              "best_params": choice["best_params"],
                              **metrics(temporal_test.cnt, pred)})

temporal_baseline_predictions = {
    "Naive lag-1": temporal_test.lag_1_cnt.to_numpy(),
    "Naive lag-7": temporal_test.lag_7_cnt.to_numpy(),
    "Naive rolling-7": temporal_test.roll_7_cnt.to_numpy(),
    "2011 constant mean": np.full(len(temporal_test), temporal_train.cnt.mean())
}
for name, pred in temporal_baseline_predictions.items():
    temporal_predictions[name] = pred
    temporal_rows_2012.append({"model": name, "feature_set": "causal baseline",
                              "best_params": "default", **metrics(temporal_test

temporal_holdout_table = (pd.DataFrame(temporal_rows_2012)
                           .sort_values("MAE").reset_index(drop=True))
temporal_holdout_table.to_csv(OUTPUT_DIR / "temporal_holdout_metrics_v5.csv", index=False)
temporal_selected_metrics = temporal_holdout_table.loc[
    temporal_holdout_table.model.eq(temporal_winner_name)
].iloc[0]
rolling_metrics = temporal_holdout_table.loc[
    temporal_holdout_table.model.eq("Naive rolling-7")
].iloc[0]
temporal_advantage = 1 - temporal_selected_metrics.MAE / rolling_metrics.MAE
print("Table 12 – Frozen temporal candidates and baselines on 2012")
display(temporal_holdout_table.round({"MAE": 1, "MSE": 1, "RMSE": 1, "R2": 3}))
```

Table 12 – Frozen temporal candidates and baselines on 2012

|   | model               | feature_set                          | best_params                        | MAE    | MSE       | RMSE   | R2     |
|---|---------------------|--------------------------------------|------------------------------------|--------|-----------|--------|--------|
| 0 | Naive rolling-7     | causal baseline                      | default                            | 847.8  | 1388485.8 | 1178.3 | 0.565  |
| 1 | Naive lag-1         | causal baseline                      | default                            | 870.2  | 1553433.8 | 1246.4 | 0.513  |
| 2 | Linear Regression   | core + interactions + autoregressive | default                            | 1047.4 | 1449639.0 | 1204.0 | 0.546  |
| 3 | Naive lag-7         | causal baseline                      | default                            | 1110.4 | 2439375.1 | 1561.8 | 0.235  |
| 4 | Random Forest       | core + interactions + autoregressive | depth=None; leaf=1; estimators=400 | 1673.3 | 3462504.7 | 1860.8 | -0.085 |
| 5 | Gradient Boosting   | core                                 | lr=0.1; depth=2; estimators=400    | 2040.5 | 4834520.0 | 2198.8 | -0.515 |
| 6 | K-Nearest Neighbors | core                                 | neighbors=3; weights=distance      | 2077.1 | 5132081.4 | 2265.4 | -0.609 |
| 7 | 2011 constant mean  | causal baseline                      | default                            | 2460.2 | 7829760.2 | 2798.2 | -1.454 |

**Temporal result.** Linear Regression with interactions and causal autoregressive inputs was frozen from 2011 at CV MAE 602.5 and CV RMSE 763.2. Adding elapsed-time trend raised CV MAE to 665.8, so the trend was excluded before 2012 evaluation. This contrast with the primary result shows that elapsed time helps interpolate within a mixed-year domain but does not establish a transferable growth law.

In 2012, the frozen Linear Regression reached MAE 1,047.4, RMSE 1,204.0 and R-squared 0.546. Rolling-7 reached MAE 847.8, RMSE 1,178.3 and R-squared 0.565, leaving the selected model 23.5% worse on MAE. Random Forest scored MAE 1,673.3, while the 2011-selected Gradient Boosting and KNN configurations exceeded 2,000. Forward temporal robustness therefore fails.

A rolling one-day-ahead test can update lags after every observed day. A forecast issued for all of 2012 on 1 January could not use those later counts and would be a different experiment with a longer horizon and recursively unavailable inputs.

## 4.3 Worked example

Aggregate error statistics hide what a prediction actually is. Table 13 takes six holdout dates that both experiments scored, and shows the same day estimated four ways: by the frozen primary model, by the constant-mean reference it must beat, by the frozen temporal model, and by the rolling seven-day rule it must beat. Reading one row left to right is the shortest description of this entire report.

```

In [18]: # Table 13 – the same dates scored by every approach compared in this report.
primary_view = primary_holdout.assign(
    primary_prediction=primary_winner.predict(primary_holdout[primary_winner_col],
    mean_reference=primary_train.cnt.mean(),
)
temporal_view = temporal_test[["dteday"]].assign(
    temporal_prediction=temporal_predictions[temporal_winner_name],
    rolling7_reference=temporal_predictions["Naive rolling-7"],
)
shared = primary_view.merge(temporal_view, on="dteday", how="inner").sort_values

# Six evenly spaced dates across the shared period, so the sample is reproducibl
positions = np.linspace(0, len(shared) - 1, 6).round().astype(int)
sample = shared.iloc[positions]

season_lbl = {1: "winter", 2: "spring", 3: "summer", 4: "autumn"}
weather_lbl = {1: "clear", 2: "mist", 3: "light rain/snow", 4: "severe"}
worked = pd.DataFrame({
    "date": sample.dteday.dt.strftime("%Y-%m-%d"),
    "season": sample.season.map(season_lbl),
    "day": sample.dteday.dt.strftime("%a"),
    "weather": sample.weathersit.map(weather_lbl),
    "temp_C": (sample.temp * 41).round(1),
    "humidity_pct": (sample.hum * 100).round(0),
    "ACTUAL": sample.cnt.astype(int),
    f"{primary_winner_name} (primary)": sample.primary_prediction.round(0).astyp
    "mean reference": sample.mean_reference.round(0).astype(int),
    f"{temporal_winner_name} (temporal)": sample.temporal_prediction.round(0).as
    "rolling-7 reference": sample.rolling7_reference.round(0).astype(int),
}).reset_index(drop=True)
worked.to_csv(OUTPUT_DIR / "worked_example_v5.csv", index=False)

print("Table 13 – Six shared dates scored by every approach")
print("Each row is one day: inputs on the left, the realised count, then four es
display(worked)

for label, column in [(f"{primary_winner_name} (primary)", "primary_prediction")
    ("mean reference", "mean_reference"),
    (f"{temporal_winner_name} (temporal)", "temporal_predictio
    ("rolling-7 reference", "rolling7_reference")]:
    print(f" mean absolute error on these six dates, {label:34s}: "
        f"{np.abs(sample.cnt - sample[column]).mean():7.0f}")
print(f"\nShared dates available: {len(shared)} (2012 days that are also primary

```

Table 13 – Six shared dates scored by every approach

Each row is one day: inputs on the left, the realised count, then four estimates of it.

|   | date       | season | day | weather | temp_C | humidity_pct | ACTUAL | Gradient Boosting (primary) | mean reference | Linear Regression (temporal) |
|---|------------|--------|-----|---------|--------|--------------|--------|-----------------------------|----------------|------------------------------|
| 0 | 2012-01-04 | winter | Wed | mist    | 4.4    | 41.0         | 2368   | 2405                        | 4575           | 1135                         |
| 1 | 2012-03-19 | winter | Mon | clear   | 22.3   | 73.0         | 6153   | 6596                        | 4575           | 5114                         |
| 2 | 2012-06-07 | spring | Thu | clear   | 24.7   | 57.0         | 7494   | 7373                        | 4575           | 6142                         |
| 3 | 2012-08-03 | summer | Fri | mist    | 31.4   | 64.0         | 7175   | 6422                        | 4575           | 5495                         |
| 4 | 2012-10-12 | autumn | Fri | clear   | 17.9   | 54.0         | 7282   | 7453                        | 4575           | 5769                         |
| 5 | 2012-12-30 | winter | Sun | clear   | 10.5   | 48.0         | 1796   | 1557                        | 4575           | 1579                         |

mean absolute error on these six dates, Gradient Boosting (primary) : 294  
mean absolute error on these six dates, mean reference : 2465  
mean absolute error on these six dates, Linear Regression (temporal) : 1173  
mean absolute error on these six dates, rolling-7 reference : 299

Shared dates available: 80 (2012 days that are also primary holdout days).

Six days cannot establish anything, and the criteria in Section 5 are computed across the full 183-day holdout and the full 366-day 2012 period for that reason. What the rows do show is the shape of each approach.

The primary model tracks the realised count in both directions, from a 2,368-rental January day to a 7,494-rental June day. The constant reference cannot move: it answers 4,575 to every question, which is why its error grows with distance from the annual mean. The temporal model follows the direction of change but sits below the realised count on the growth days, because it was fitted on a smaller 2011 system and has no mechanism for the expansion that followed. The rolling seven-day rule carries no model at all and stays close, because each new observed day re-anchors the next estimate.

The primary and rolling-7 columns must not be read as a head-to-head contest. They belong to different comparisons in Table 1: the primary model was fitted on days drawn from both years, including days adjacent to these, and answers a conditional question; the rolling rule answers a forward one and needs no training. Their apparent similarity on six shared dates is a property of this sample, not a finding. The comparison that matters for each is in Tables 7 and 12 respectively.

## 5. Evaluation

### 5.1 Criteria summary and paired temporal description

The primary assessment criterion and secondary robustness criterion are reported separately. The temporal comparison also retains descriptive pairing: median absolute error and the share of identical 2012 dates won by each method. These summaries do not require an independence assumption. No inferential significance claim is made because adjacent daily errors are temporally related.

```
In [19]: primary_pass = primary_improvement >= 0.40
temporal_pass = temporal_advantage >= 0.05
evaluation_summary = pd.DataFrame([
    {"evaluation": "Conditional demand estimation",
     "criterion": ">= 40% MAE improvement over training-mean baseline",
     "measured": f"{primary_improvement*100:.1f}% improvement",
     "result": "Meets" if primary_pass else "Does not meet"},
    {"evaluation": "Forward temporal robustness",
     "criterion": ">= 5% MAE advantage over rolling-7",
     "measured": f"{temporal_advantage*100:.1f}% advantage",
     "result": "Passes" if temporal_pass else "Fails"},
])
evaluation_summary.to_csv(OUTPUT_DIR / "evaluation_summary_v5.csv", index=False)
print("Table 14 - Evaluation summary")
display(evaluation_summary)

y2012 = temporal_test.cnt.to_numpy()
selected_temporal_pred = temporal_predictions[temporal_winner_name]
rolling_pred = temporal_predictions["Naive rolling-7"]
selected_error = np.abs(y2012 - selected_temporal_pred)
rolling_error = np.abs(y2012 - rolling_pred)
ties = selected_error == rolling_error
paired_summary = pd.DataFrame([
    {"series": temporal_winner_name, "median_absolute_error": np.median(selected_error),
     "days_closer": int((selected_error < rolling_error).sum()),
     "win_rate_excluding_ties_pct": 100 * (selected_error < rolling_error).sum() / len(selected_error)},
    {"series": "Naive rolling-7", "median_absolute_error": np.median(rolling_error),
     "days_closer": int((rolling_error < selected_error).sum()),
     "win_rate_excluding_ties_pct": 100 * (rolling_error < selected_error).sum() / len(rolling_error)},
])
paired_summary["tied_dates"] = int(ties.sum())
paired_summary.to_csv(OUTPUT_DIR / "paired_temporal_comparison_v5.csv", index=False)
print("Table 15 - Paired descriptive comparison on identical 2012 dates")
display(paired_summary.round(1))
```

Table 14 - Evaluation summary

|   | evaluation                    | criterion                                         | measured          | result |
|---|-------------------------------|---------------------------------------------------|-------------------|--------|
| 0 | Conditional demand estimation | >= 40% MAE improvement over training-mean base... | 73.9% improvement | Meets  |
| 1 | Forward temporal robustness   | >= 5% MAE advantage over rolling-7                | -23.5% advantage  | Fails  |

Table 15 - Paired descriptive comparison on identical 2012 dates



|   | series            | median_absolute_error | days_closer | win_rate_excluding_ties_pct | tied_dates |
|---|-------------------|-----------------------|-------------|-----------------------------|------------|
| 0 | Linear Regression | 1045.2                | 116         | 31.7                        | 0          |
| 1 | Naive rolling-7   | 636.1                 | 250         | 68.3                        | 0          |

```
In [20]: train_ceiling = int(temporal_train.cnt.max())
ceiling_rows = []
for name in list(GRIDS) + ["Naive rolling-7"]:
    pred = temporal_predictions[name]
    ceiling_rows.append({
        "model": name,
        "largest_2012_prediction": float(np.max(pred)),
        "2011_training_ceiling": train_ceiling,
        "2012_actual_maximum": int(temporal_test.cnt.max()),
        "mean_residual_actual_minus_prediction": float(np.mean(y2012 - pred)),
    })
extrapolation_ceiling = pd.DataFrame(ceiling_rows)
extrapolation_ceiling.to_csv(OUTPUT_DIR / "extrapolation_ceiling_v5.csv", index=
print("Table 16 - Extrapolation ceiling analysis")
display(extrapolation_ceiling.round(1))

fig, ax = plt.subplots(1, 2, figsize=(14, 5))
ax[0].plot(temporal_test.dteday, y2012, lw=.9, color="#20639B", label="actual 20
ax[0].plot(temporal_test.dteday, selected_temporal_pred, lw=.9, color="#ED553B",
            label=f"frozen {temporal_winner_name}")
ax[0].plot(temporal_test.dteday, rolling_pred, lw=.9, color="#3CAEA3", label="ro
ax[0].axhline(train_ceiling, color="black", ls="--", lw=1, label=f"2011 ceiling
ax[0].set(title="Figure 10 - Rolling day-ahead predictions in 2012", xlabel="dat
ax[0].legend(fontsize=8)
ax[1].scatter(selected_temporal_pred, y2012-selected_temporal_pred, s=14, alpha=
ax[1].axhline(0, color="black", ls="--")
ax[1].set(title=f"Figure 11 - {temporal_winner_name} temporal residuals",
            xlabel="predicted cnt", ylabel="actual - predicted")
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_temporal_robustness.png", dpi=160)
```

Table 16 - Extrapolation ceiling analysis

|   | model               | largest_2012_prediction | 2011_training_ceiling | 2012_actual_maximum | mean_residua |
|---|---------------------|-------------------------|-----------------------|---------------------|--------------|
| 0 | Linear Regression   | 8247.5                  | 6043                  | 8714                |              |
| 1 | K-Nearest Neighbors | 5540.7                  | 6043                  | 8714                |              |
| 2 | Random Forest       | 5567.2                  | 6043                  | 8714                |              |
| 3 | Gradient Boosting   | 5490.7                  | 6043                  | 8714                |              |
| 4 | Naive rolling-7     | 7988.0                  | 6043                  | 8714                |              |



Tree ensembles predict averages within terminal regions and therefore have an extrapolation ceiling tied to observed training targets. The 2011 maximum was 6,043 rentals, while 2012 reached 8,714. Random Forest's largest prediction was 5,567 and Gradient Boosting's was 5,491; their mean under-predictions were 1,568 and 1,995 rentals per day. KNN also stayed below the training ceiling.

Linear Regression extended to 8,248 and reduced mean under-prediction to 832, explaining why it led the 2011-selected ML candidates. Rolling-7 reached 7,988 and had a mean residual of -2.5 because each new observation re-anchored the next prediction. Conditional accuracy within a mixed-year sample therefore did not transfer to the later growth period.

## 5.2 Explaining the frozen primary model

Permutation importance is model-agnostic and measures the increase in holdout MAE after one raw input is shuffled. It is computed only after model selection and does not change the winner. Contributions describe predictive associations in this dataset, not causal effects. TreeSHAP (Lundberg & Lee, 2017) is added only when the frozen primary family is Random Forest or Gradient Boosting.


```
In [21]: perm = permutation_importance(
    primary_winner, primary_holdout[primary_winner_columns], primary_holdout.cnt,
    scoring="neg_mean_absolute_error", n_repeats=30, random_state=RANDOM_STATE,
)
permutation_table = (pd.DataFrame({
    "feature": primary_winner_columns,
    "MAE_increase_mean": perm.importances_mean,
    "MAE_increase_std": perm.importances_std,
}).sort_values("MAE_increase_mean", ascending=False).reset_index(drop=True))
permutation_table.to_csv(OUTPUT_DIR / "permutation_importance_v5.csv", index=False)
print("Table 17 - Model-agnostic permutation importance on the primary holdout")
display(permutation_table.head(12).round(2))

top_perm = permutation_table.head(12).sort_values("MAE_increase_mean")
fig, ax = plt.subplots(figsize=(9, 5.5))
ax.barh(top_perm.feature, top_perm.MAE_increase_mean, xerr=top_perm.MAE_increase_std,
        color="#20639B", alpha=.85)
ax.set(title=f"Figure 12 - Permutation importance: {primary_winner_name}",
        xlabel="increase in holdout MAE after permutation", ylabel="")
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_permutation_importance.png", dpi=100)
```

Table 17 - Model-agnostic permutation importance on the primary holdout

|    | feature          | MAE_increase_mean | MAE_increase_std |
|----|------------------|-------------------|------------------|
| 0  | days_since_start | 950.28            | 53.08            |
| 1  | atemp            | 409.77            | 41.65            |
| 2  | hum              | 144.47            | 18.57            |
| 3  | temp_sq          | 110.50            | 16.63            |
| 4  | mnth             | 81.71             | 15.54            |
| 5  | windspeed        | 76.41             | 15.94            |
| 6  | weathersit       | 69.47             | 9.73             |
| 7  | workingday       | 34.08             | 9.05             |
| 8  | atemp_hum        | 29.41             | 6.80             |
| 9  | temp             | 24.92             | 8.24             |
| 10 | season           | 24.35             | 5.03             |
| 11 | weekday          | 23.60             | 7.73             |


Figure 12 - Permutation importance: Gradient Boosting





```
In [22]: shap_table = None
if primary_winner_name in {"Random Forest", "Gradient Boosting"}:
    import shap
    fitted_prep = primary_winner.named_steps["prep"]
    transformed = fitted_prep.transform(primary_holdout[primary_winner_columns])
    if hasattr(transformed, "toarray"):
        transformed = transformed.toarray()
    transformed_names = fitted_prep.get_feature_names_out()
    explainer = shap.TreeExplainer(primary_winner.named_steps["model"])
    shap_values = explainer.shap_values(transformed, check_additivity=True)
    shap_table = (pd.DataFrame({"feature": transformed_names,
                               "mean_absolute_SHAP": np.abs(shap_values).mean(
                                   axis=1).sort_values("mean_absolute_SHAP", ascending=False).reset_index(
                                   level=1).feature_name}),
                  shap_table.to_csv(OUTPUT_DIR / "shap_global_importance_v5.csv", index=False))
    print("Table 18 - TreeSHAP global importance for the frozen primary model")
    display(shap_table.head(12).round(2))
    shap.summary_plot(shap_values, transformed, feature_names=transformed_names,
                      max_display=12, show=False)
    plt.title(f"Figure 13 - TreeSHAP summary: {primary_winner_name}")
    plt.tight_layout(); plt.savefig(FIG_DIR / "v5_shap_summary.png", dpi=160, bb
else:
    print("TreeSHAP omitted: the frozen primary winner is", primary_winner_name)
```

Table 18 - TreeSHAP global importance for the frozen primary model

|    | feature                      | mean_absolute_SHAP |
|----|------------------------------|--------------------|
| 0  | continuous__days_since_start | 1018.06            |
| 1  | continuous__atemp            | 586.20             |
| 2  | continuous__hum              | 234.58             |
| 3  | continuous__temp_sq          | 201.12             |
| 4  | continuous__windspeed        | 136.82             |
| 5  | continuous__temp             | 119.90             |
| 6  | nominal__weathersit_2        | 97.63              |
| 7  | binary__workingday           | 87.97              |
| 8  | nominal__weekday_6           | 69.75              |
| 9  | nominal__mnth_8              | 55.72              |
| 10 | nominal__weathersit_3        | 55.56              |
| 11 | continuous__atemp_hum        | 52.08              |



```
In [23]: primary_pred = primary_winner.predict(primary_holdout[primary_winner_columns])
fig, ax = plt.subplots(1, 2, figsize=(13, 5))
ax[0].scatter(primary_holdout.cnt, primary_pred, alpha=.45, s=18, color="#20639B")
lims = [min(primary_holdout.cnt.min(), primary_pred.min()),
        max(primary_holdout.cnt.max(), primary_pred.max())]
ax[0].plot(lims, lims, "k--")
ax[0].set(title="Figure 14 - Primary holdout: predicted vs actual",
          xlabel="actual cnt", ylabel="predicted cnt")
ax[1].scatter(primary_pred, primary_holdout.cnt-primary_pred, alpha=.45, s=18, color="black")
ax[1].axhline(0, color="black", ls="--")
ax[1].set(title="Figure 15 - Primary holdout residuals",
          xlabel="predicted cnt", ylabel="actual - predicted")
plt.tight_layout(); plt.savefig(FIG_DIR / "v5_primary_diagnostics.png", dpi=160)
```



## 5.3 Interpretation

Permutation confirms that elapsed time is the dominant predictive association in the primary holdout: shuffling it increases MAE by about 950 rentals. Feels-like temperature follows at 410, then humidity at 144 and squared temperature at 111. TreeSHAP produces the same leading order, with mean absolute contributions of about 1,018, 586, 235 and 201 rentals. These values explain predictions across the observed sample; they do not imply that advancing time causes demand or that weather effects are causal.

The primary result is MAE 433.9, RMSE 630.6 and R-squared 0.897: 73.9% better than the baseline and 10.1% of mean holdout demand. Temporally, Linear Regression lost to rolling-7 by 23.5%; median error was 1,045 versus 636, with wins on 116 versus 250 dates. This paired description makes no independence claim.

## 6. Lessons Learned

### 6.1 What went well

The method built during Assessments 1 and 2 carried straight into this project and shortened the setup considerably. Declaring criteria before modelling, keeping a holdout untouched until selection finishes, running an ablation to attribute a metric change to a cause, and operating the notebook as a submission artefact were all decided in earlier assessments, so this one reached a working pipeline much faster. The remaining effort went into the questions specific to this dataset instead of into process.

The design also held up under its own checks. All 731 daily units remain available for the primary question, the holdout is isolated from every selection decision, and four families compete across three exogenous feature sets. The hourly file validates every daily total and exposes both intraday peaks and panel imbalance.

### 6.2 Challenges

Understanding the data took longer than fitting the models. Two files describe the same system at different granularities, `casual` and `registered` sum to the target and cannot be predictors, and one humidity reading is physically impossible. Diagnosing that reading required cross-checking the hourly file, and the correction then had to preserve sample size and stay inside the pipeline folds.

The concept that took longest to internalise is the difference between interpolation and extrapolation. A tree fits both years comfortably when they are shuffled together, and the same tree cannot reach 2012 demand when trained on 2011 alone, because a terminal region can only average values it has already seen. That is a property of the model class rather than a tuning failure, and it explains why the simplest candidate transferred best.

### 6.3 What can be improved

The strongest limitations are evidential. More years, archived weather forecasts and repeated forward windows would test whether the temporal result reflects this transition or a permanent property. Station-level planning would additionally need identifiers, dock capacity, trip origins and destinations, transit connections and local demand, which is why the ITDP and NACTO guidance points beyond this dataset. Category-4 weather remains outside the supported domain.

Methodologically, trend-seasonality decomposition, differenced targets and forecast-aware inputs are the natural next comparisons.

The extension I would most like to build is a small service that turns the frozen primary model into something usable: send it a date, its season and a weather forecast, and

receive an estimated daily volume with the supported-domain guard applied.  
Appendix C sketches that idea, including the parts this study shows it must not claim.

---

## Appendix A - Glossary

| Term                          | Meaning in this report                                                                          |
|-------------------------------|-------------------------------------------------------------------------------------------------|
| <b>MAE</b>                    | Mean absolute error: average miss in rentals per day.                                           |
| <b>MSE / RMSE</b>             | Squared error and its root; large misses receive greater weight.                                |
| <b>R-squared</b>              | Variance explained relative to the evaluated sample mean; it may be negative.                   |
| <b>Holdout</b>                | Data isolated from model, feature and hyperparameter selection and used once for final scoring. |
| <b>Interpolation</b>          | Prediction within a domain represented in training.                                             |
| <b>Extrapolation</b>          | Prediction beyond the observed training range or period.                                        |
| <b>Autoregressive feature</b> | Input constructed from demand observed before the target day.                                   |
| <b>TimeSeriesSplit</b>        | Expanding-window cross-validation that preserves temporal order.                                |
| <b>Permutation importance</b> | Model-agnostic holdout score degradation after shuffling an input.                              |
| <b>SHAP</b>                   | Additive prediction attribution used here only for a frozen tree winner.                        |

## Appendix B - Reproducibility and integrity evidence

```
In [24]: import sys
import sklearn
import matplotlib

assert len(primary_df) == 731
assert len(temporal_df) == 724
assert selection_indices.isdisjoint(set(primary_holdout_idx))
assert temporal_train.yr.eq(0).all()
assert temporal_train.dteday.max() < temporal_test.dteday.min()
assert all("casual" not in cols and "registered" not in cols
          for cols in list(PRIMARY_FEATURE_SETS.values()) + list(TEMPORAL_FEATU
assert all(c not in cols for cols in PRIMARY_FEATURE_SETS.values() for c in AUTO
assert cross_file.difference.eq(0).all() and len(cross_file) == 731
assert len(primary_holdout_table) == 5 and len(temporal_holdout_table) == 8

integrity = pd.DataFrame([
    ("Primary rows retained", 731, len(primary_df), "Pass"),
    ("Temporal rows after causal warm-up", 724, len(temporal_df), "Pass"),
    ("Primary holdout rows excluded from selection", 183, len(primary_holdout_id
    ("Temporal selection years", "2011 only", str(temporal_train.dteday.dt.year.
    ("Hourly/daily dates reconciled", 731, int(cross_file.difference.eq(0).sum()
    ("Dates with incomplete hourly panels", 76, incomplete_dates, "Pass"),
    ("Omitted zero-demand hourly rows", 165, omitted_hour_rows, "Pass"),
    ("Leakage components in predictor lists", 0, 0, "Pass"),
], columns=["check", "expected", "observed", "result"])
integrity.to_csv(OUTPUT_DIR / "integrity_checks_v5.csv", index=False)
display(integrity)
print("Python:", sys.version.split()[0], "| pandas:", pd.__version__,
      "| scikit-learn:", sklearn.__version__, "| matplotlib:", matplotlib.__vers
print("RANDOM_STATE:", RANDOM_STATE)
```

|   | check                                        | expected  | observed | result |
|---|----------------------------------------------|-----------|----------|--------|
| 0 | Primary rows retained                        | 731       | 731      | Pass   |
| 1 | Temporal rows after causal warm-up           | 724       | 724      | Pass   |
| 2 | Primary holdout rows excluded from selection | 183       | 183      | Pass   |
| 3 | Temporal selection years                     | 2011 only | [2011]   | Pass   |
| 4 | Hourly/daily dates reconciled                | 731       | 731      | Pass   |
| 5 | Dates with incomplete hourly panels          | 76        | 76       | Pass   |
| 6 | Omitted zero-demand hourly rows              | 165       | 165      | Pass   |
| 7 | Leakage components in predictor lists        | 0         | 0        | Pass   |

Python: 3.14.3 | pandas: 3.0.3 | scikit-learn: 1.9.0 | matplotlib: 3.11.0  
RANDOM\_STATE: 42

## Appendix C - Proposed extension: a demand estimation service

This appendix sketches the engineering extension named in Section 6.3. It is a proposal, not a deployed system, and the evidence in Section 5 constrains what it may offer.

### C.1 What it would do

A caller sends a date, its season and a weather forecast; the service returns an estimated daily rental volume from the frozen primary model, together with the supported input domain.

```
POST /estimate
{ "date": "2013-04-18", "season": 2, "weathersit": 1,
  "temp": 0.58, "atemp": 0.55, "hum": 0.61, "windspeed": 0.18 }

-> { "estimated_daily_rentals": 5120,
    "expected_absolute_error": 434,
    "basis": "conditional estimate, observed 2011-2012 distribution",
    "supported_weathersit": [1, 2, 3] }
```

Returning the expected error alongside the point estimate matters more than the estimate itself. A planner who receives 5,120 without 434 will read a precision the model does not have.

### C.2 Proposed layers

| Layer          | Technology                    | Purpose                                                                                            |
|----------------|-------------------------------|----------------------------------------------------------------------------------------------------|
| Model artefact | scikit-learn pipeline, joblib | The frozen Gradient Boosting configuration and its preprocessing, serialised together              |
| Contract       | JSON                          | Feature order, supported <code>weathersit</code> categories, holdout metrics, source notebook hash |
| API            | FastAPI, Pydantic             | Typed request validation and the domain guard                                                      |
| Interface      | Streamlit                     | A form for planners who will not call an API                                                       |
| Monitoring     | Scheduled job                 | Daily scoring of yesterday's estimate against realised demand                                      |

The pattern follows my Sommelier API project from Assessments 1 and 2, where the same separation between a framework-agnostic model core and thin serving surfaces is already implemented.

### C.3 What this study forbids it from claiming

The temporal check is the reason this appendix is short about forecasting.

- The service may answer **conditional** questions: given these conditions, what does the 2011-2012 relationship imply? That is the question the primary experiment evaluated.
- It may **not** be presented as a next-day forecaster. Section 4.2 measured the frozen model at 23.5% worse than a rolling seven-day mean on 2012. A

forecasting endpoint would either serve the rolling mean or wait for new forward evidence.

- Any request outside `weathersit` 1 to 3 must be rejected for human review rather than scored, because category 4 never appears in training.

## C.4 Geolocation and alerts

Two ideas worth recording, with their preconditions.

**Station-level geolocation.** Plotting expected demand onto a map through a mapping API is straightforward as an interface, and unsupported as an inference. This dataset carries no station identifier, coordinate, dock capacity or trip origin, so nothing in this notebook can attribute system demand to a location. Delivering it honestly means acquiring station-level data first; ITDP and NACTO describe the network, spacing and capacity variables that work would need (ITDP, 2018; NACTO, 2016).

**Custom alerts.** Two alert types are supported by evidence already collected. A *domain alert* fires when a request falls outside the supported weather categories. A *drift alert* fires when the rolling error of the deployed estimate exceeds its holdout MAE over a defined window, which is the mechanism that would have caught the temporal failure in production rather than in a notebook. A demand-threshold alert for staffing would be useful but needs a stakeholder-defined threshold that this project does not have.

---

## Academic Integrity Declaration

I declare that this submission is my own work. All sources of information, ideas and code have been acknowledged. The dataset is publicly available from the UCI Machine Learning Repository. The analysis, modelling decisions, interpretation and written commentary are my own.

## Statement of Acknowledgement

Analysis was performed in Python with pandas, NumPy, scikit-learn, Matplotlib, Seaborn and SHAP.

I acknowledge that I have used the following AI tool(s) in the creation of this report:

- Anthropic Claude Opus 5
- OpenAI ChatGPT Codex 5.6

Both tools were used to assist with understanding ML concepts, challenging the experimental design, auditing empirical claims against executed outputs, structuring the technical pipeline, improving clarity of academic language, and supporting APA 7th referencing conventions.

Prompt examples:

1. "Explain why a random split over a two-year time series supports a conditional demand estimate but not a next-day forecast, and how to run both experiments so that neither contaminates the other."
2. "Show why a tree ensemble trained on 2011 cannot predict 2012 demand above its training maximum, and propose a way to measure that ceiling instead of asserting it."
3. "Build a causal seven-day rolling baseline for a day-ahead comparison using shift before rolling, and write the assertion that proves the window never includes the target day."

I confirm that the use of these tools has been in accordance with the Torrens University Australia Academic Integrity Policy and TUA, Think and MDS's Position Paper on the Use of AI. I confirm that the final output is authored by me and represents my own critical thinking, analysis, and synthesis of sources. I take full responsibility for the final content of this report.

## References

Capital Bikeshare. (n.d.). *How it works*. <https://capitalbikeshare.com/how-it-works>

Chapman, P., Clinton, J., Kerber, R., Khabaza, T., Reinartz, T., Shearer, C., & Wirth, R. (2000). *CRISP-DM 1.0: Step-by-step data mining guide*. SPSS Inc.

Fanaee-T, H., & Gama, J. (2014). Event labeling combining ensemble detectors and background knowledge. *Progress in Artificial Intelligence*, 2(2-3), 113-127.

<https://doi.org/10.1007/s13748-013-0040-3>

Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29(5), 1189-1232. <https://doi.org/10.1214/aos/1013203451>

Institute for Transportation & Development Policy. (2018). *The bikeshare planning guide*. <https://itdp.org/publication/the-bike-share-planning-guide/>

Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765-4774.

National Association of City Transportation Officials. (2016). *Bike share station siting guide*. <https://nacto.org/publication/bike-share-station-siting-guide/>

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.

University of California, Irvine. (n.d.). *Bike Sharing Dataset*. UCI Machine Learning Repository. <https://archive.ics.uci.edu/dataset/275/bike+sharing+dataset>