

Predicting Daily Demand for Capital Bikeshare

Regression Model Comparison and Temporal Robustness Testing

MLN601 Assessment 3 - CRISP-DM

Design and Creative Technologies, Torrens University

- **Student:** Luis Guilherme de Barros Andrade Faria - A00187785
- **Subject:** Machine Learning (MLN601)
- **Lecturer:** Dr. Kamran Shaukat
- **Assessment:** 3
- **Date:** August 2026

Field	Scope
Dataset	UCI Bike Sharing, Capital Bikeshare daily and hourly files, Washington DC, 2011-2012
Target	cnt : system-wide daily rentals
Primary question	Conditional daily-demand estimation from calendar attributes and expected weather
Secondary check	Rolling day-ahead temporal robustness, trained in 2011 and evaluated in 2012
Method	Regression comparison within CRISP-DM (Chapman et al., 2000)

1. Business Understanding

1.1 Business objective

Capital Bikeshare is a docked bicycle-sharing system (Capital Bikeshare, n.d.). This project asks:

Given calendar attributes and expected weather conditions, how accurately can regression models estimate Capital Bikeshare's system-wide daily rental demand?

Daily demand is total network rentals in one day. An accurate estimate could support fleet-capacity, staffing and maintenance planning. Station inventory and rebalancing are outside scope because the CSVs lack station identifiers, locations and capacity; station-level data could extend the work.

Lime in Sydney is a contemporary dockless analogy that motivates my interest and makes the assessment closer to my reality. Although its operations differ, it raises a parallel question: how could calendar and weather patterns inform fleet planning in a local system? It motivates the work but is not model evidence.

1.2 Data and prediction scope

Each row represents one system-wide day, with `cnt` as total rentals. Calendar attributes are known in advance. Observed weather proxies expected weather; real use would require forecasts. The hourly file supports exploration without changing the daily target.

The primary random 75/25 experiment estimates demand for days resembling 2011-2012 using calendar and weather without past demand. The secondary experiment tests transfer through time: selection uses 2011, then counts through day D plus next-day conditions predict D+1 in 2012. Forecasting all of 2012 on 1 January is separate because later counts would be unavailable.

This report therefore runs three comparisons. The first two answer the assessment question; the third tests whether that answer transfers forward in time.

Table 1 - What this report compares

#	Comparison	Split	Compared against	Question answered	Where
1	Model selection	Random 75/25, training partition only	The other 11 family and feature-set configurations	Which configuration predicts most accurately?	Table 6
2	Value of modelling	Random 75/25, holdout opened once	A constant training-mean prediction	Does the selected model beat not modelling at all?	Tables 7 and 8

#	Comparison	Split	Compared against	Question answered	Where
3	Forward transfer	Time-ordered: trained in 2011, scored on 2012	Naive lag-1, lag-7 and rolling seven-day rules	Does the approach still work on a period it has not seen?	Table 12

Comparisons 1 and 2 address the brief. Comparison 3 is an additional check that limits how the result may be used.

1.3 Evaluation criteria

MAE is primary because it reports the average miss in rentals per day. RMSE gives more weight to large errors, while R-squared shows explained variation relative to predicting the sample mean.

The training-mean baseline tests whether predictors add value; the primary model must improve MAE by 40%. Rolling seven-day demand is the temporal reference because it adapts to recent demand; the frozen ML model must beat it by 5% on identical 2012 dates.

Model, feature and hyperparameter choices use training-only cross-validation. Lowest MAE leads; configurations within 5% are separated by RMSE and then simplicity. The holdout is used only after selection.

These study-defined criteria came from earlier exploratory versions and are not stakeholder-validated service levels.

1.4 Assumptions and limitations

- Observed weather omits forecast error, making reported day-ahead accuracy optimistic.
- Two years provide one year-on-year transition, weak evidence for persistent growth.
- Weather category 4 is absent; an input guard must restrict use to supported categories 1-3.
- System totals cannot answer station questions. The Institute for Transportation and Development Policy (ITDP) publishes bikeshare planning guidance; the National Association of City Transportation Officials (NACTO) publishes network and station-siting guidance. They explain why future station planning also needs identifiers, capacity, trip flows, transit connections and local demand (ITDP, 2018; NACTO, 2016).

2. Data Understanding

2.1 Source, loading and granularity

The UCI Bike Sharing dataset supplies `day.csv` and `hour.csv` for Capital Bikeshare in 2011-2012 (Fanaee-T & Gama, 2014; University of California, Irvine, n.d.). The daily file is the modelling source. The hourly file is loaded separately for intraday EDA and cross-file validation.

```
In [1]: import os
import warnings
from pathlib import Path

os.environ["PYTHONWARNINGS"] = "ignore"
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.base import clone
from sklearn.dummy import DummyRegressor
from sklearn.linear_model import LinearRegression
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import (
    train_test_split, GridSearchCV, KFold, TimeSeriesSplit, cross_validate
)
from sklearn.inspection import permutation_importance
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

sns.set_theme(style="whitegrid")
RANDOM_STATE = 42

NB_DIR = Path.cwd()
BASE_DIR = NB_DIR.parent if NB_DIR.name == "notebook" else NB_DIR
DATA_DIR = BASE_DIR / "dataset"
OUTPUT_DIR = BASE_DIR / "outputs"
FIG_DIR = OUTPUT_DIR / "figures"
FIG_DIR.mkdir(parents=True, exist_ok=True)

day_df = pd.read_csv(DATA_DIR / "day.csv", parse_dates=["dteday"]).sort_values("dteday")
hour_df = pd.read_csv(DATA_DIR / "hour.csv", parse_dates=["dteday"]).sort_values(["dteday",
print("day.csv :", day_df.shape, day_df.dteday.min().date(), "to", day_df.dteday.max().date())
print("hour.csv:", hour_df.shape, hour_df.dteday.min().date(), "to", hour_df.dteday.max().date())

day.csv : (731, 16) 2011-01-01 to 2012-12-31
hour.csv: (17379, 17) 2011-01-01 to 2012-12-31
```

```
In [2]: hourly_daily = hour_df.groupby("dteday", as_index=False)["cnt"].sum().rename(columns={"cnt": "hourly_sum"})
cross_file = day_df[["dteday", "cnt"]].merge(hourly_daily, on="dteday", how="outer", validate="one-to-one")
cross_file["difference"] = cross_file["hourly_sum"] - cross_file["cnt"]
assert len(cross_file) == 731
assert cross_file["difference"].eq(0).all(), "Hourly cnt does not sum to daily cnt for even dates"
hours_per_date = hour_df.groupby("dteday").size()
incomplete_dates = int(hours_per_date.lt(24).sum())
omitted_hour_rows = int((24 - hours_per_date).sum())
assert incomplete_dates == 76 and omitted_hour_rows == 165

granularity = pd.DataFrame([
    {"file": "day.csv", "rows": len(day_df), "observation": "one daily prediction unit",
     "role": "selected for modelling"},
    {"file": "hour.csv", "rows": len(hour_df), "observation": "correlated hourly observations",
     "role": "intraday understanding and cross-file validation"},
])
print("Table 2 - Granularity and role")
display(granularity)
print(f"Cross-file validation passed: hourly cnt sums exactly to daily cnt on all {len(cross_file)} dates.")
print(f"Hourly panel audit: {incomplete_dates} dates have fewer than 24 rows; "
      f"{omitted_hour_rows} zero-demand hour rows are omitted in total.")
```

Table 2 – Granularity and role

	file	rows	observation	role
0	day.csv	731	one daily prediction unit	selected for modelling
1	hour.csv	17379	correlated hourly observations	intraday understanding and cross-file validation

Cross-file validation passed: hourly cnt sums exactly to daily cnt on all 731 dates.
Hourly panel audit: 76 dates have fewer than 24 rows; 165 zero-demand hour rows are omitted in total.


`hour.csv` is unbalanced: 76 of 731 dates have fewer than 24 observations, with 165 zero-demand rows omitted. An hourly model would change granularity, weight dates unequally and omit those periods. Selecting `day.csv` gives every target day one complete observation **to avoid hourly bias**. The hourly counts still reconstruct all daily totals and remain useful for intraday interpretation.

```
In [3]: hour_profile = (hour_df.groupby(["workingday", "hr"], as_index=False)["cnt"].mean())
peaks = (hour_profile.loc[hour_profile.groupby("workingday")["cnt"].idxmax()]
        .assign(day_type=lambda x: x["workingday"].map({0: "Non-working day", 1: "Working day"}))
        .rename(columns={"hr": "peak_hour", "cnt": "mean_cnt"}))
peaks["mean_cnt"] = peaks["mean_cnt"].round(1)
print("Table 3 - Peak hourly demand")
display(peaks)
assert int(peaks.loc[peaks.day_type == "Working day", "peak_hour"].iloc[0]) == 17
assert np.isclose(peaks.loc[peaks.day_type == "Working day", "mean_cnt"].iloc[0], 525.3)
assert int(peaks.loc[peaks.day_type == "Non-working day", "peak_hour"].iloc[0]) == 13
assert np.isclose(peaks.loc[peaks.day_type == "Non-working day", "mean_cnt"].iloc[0], 372.7)

fig, ax = plt.subplots(figsize=(10.5, 4.6))
for flag, label, color in [(1, "Working day", "#20639B"), (0, "Non-working day", "#ED553B")]:
    p = hour_profile[hour_profile.workingday == flag]
    ax.plot(p.hr, p.cnt, marker="o", ms=3, lw=2, label=label, color=color)
ax.set(title="Figure 1 - Mean hourly rentals by day type", xlabel="hour of day", ylabel="mean hourly cnt")
ax.set_xticks(range(0, 24, 2)); ax.legend()
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_hourly_demand.png", dpi=160); plt.show()
```

Table 3 - Peak hourly demand

	day_type	peak_hour	mean_cnt
13	Non-working day	13	372.7
41	Working day	17	525.3



Working-day demand peaks at 17:00 (525.3 rentals), while non-working-day demand peaks at 13:00 (372.7). The profiles are consistent with commute-oriented and broader daytime use, respectively; aggregate counts do not prove trip purpose.

```
In [4]: span_days = (day_df.dteday.max() - day_df.dteday.min()).days + 1
norm_cols = ["temp", "atemp", "hum", "windspeed"]
quality = pd.DataFrame([
    ("Daily rows equal calendar span", len(day_df) == span_days, f"{len(day_df)} rows / {span_days}"),
    ("No date gaps", day_df.dteday.diff().dropna().dt.days.max() == 1, "daily continuity"),
    ("No exact duplicates", day_df.duplicated().sum() == 0, f"{day_df.duplicated().sum()} duplicates"),
    ("No source missing values", day_df.isna().sum().sum() == 0, f"{day_df.isna().sum().sum()} missing"),
    ("casual + registered = cnt", (day_df.casual + day_df.registered == day_df.cnt).all(), "all rows"),
    ("Normalised weather in [0,1]", all(day_df[c].between(0, 1).all() for c in norm_cols), "all four columns"),
    ("Humidity above zero", (day_df.hum > 0).all(), f"{int((day_df.hum == 0).sum())} zero readings"),
    ("Daily weather categories supported", set(day_df.weathersit.unique()) == {1, 2, 3}, "category 4 absent"),
], columns=["check", "pass", "evidence"])
print("Table 4 – Data quality audit")
display(quality)

zero_hum_date = day_df.loc[day_df.hum.eq(0), "dteday"].iloc[0]
zero_hour_rows = hour_df[hour_df.dteday.eq(zero_hum_date)]
print("Zero-humidity date:", zero_hum_date.date(), "| hourly rows:", len(zero_hour_rows),
      "| hourly humidity values:", sorted(zero_hour_rows.hum.unique()))

# Keep a raw frame for sensitivity analysis, then mark the implausible value missing.
raw_day_df = day_df.copy()
day_df.loc[day_df.hum.eq(0), "hum"] = np.nan
assert day_df.hum.isna().sum() == 1
```

Table 4 – Data quality audit

	check	pass	evidence
0	Daily rows equal calendar span	True	731 rows / 731 days
1	No date gaps	True	daily continuity
2	No exact duplicates	True	0 duplicates
3	No source missing values	True	0 missing
4	casual + registered = cnt	True	all rows
5	Normalised weather in [0,1]	True	all four columns
6	Humidity above zero	False	1 zero reading
7	Daily weather categories supported	True	category 4 absent

Zero-humidity date: 2011-03-10 | hourly rows: 22 | hourly humidity values: [np.float64(0.0)]

The daily series is complete and `casual + registered = cnt` holds throughout. One date reports zero humidity across 22 hourly rows. It is treated as a recording fault, changed to missing and median-imputed inside each training fold; a sensitivity check tests the decision. Plausible tail values are retained.

Weather category 4 is absent, so the supported domain is 1-3 and still requires an input guard.

2.2 Variables and exploratory analysis

Variables	Role
<code>cnt</code>	Continuous target: system-wide rentals per day
<code>season</code> , <code>yr</code> , <code>mnth</code> , <code>holiday</code> , <code>weekday</code> , <code>workingday</code>	Calendar predictors
<code>weathersit</code> , <code>temp</code> , <code>atemp</code> , <code>hum</code> , <code>windspeed</code>	Weather predictors
<code>dteday</code>	Ordering and elapsed-time construction; not passed directly to a model
<code>casual</code> , <code>registered</code>	Descriptive plots only; excluded because they sum exactly to <code>cnt</code>

`casual` and `registered` describe user groups, not independent predictors. Registered users may include commuters, but membership status does not prove trip purpose.


```

In [5]: season_lbl = {1: "winter", 2: "spring", 3: "summer", 4: "autumn"}
weather_lbl = {1: "clear", 2: "mist", 3: "light rain/snow", 4: "heavy rain/snow"}
day_df["season_name"] = day_df.season.map(season_lbl)
day_df["weather_name"] = day_df.weathersit.map(weather_lbl)
day_df["year_name"] = day_df.yr.map({0: "2011", 1: "2012"})

fig, ax = plt.subplots(2, 2, figsize=(13, 9))
sns.histplot(day_df.cnt, bins=30, kde=True, ax=ax[0, 0], color="#20639B")
ax[0, 0].set(title="Daily-rental distribution", xlabel="daily cnt")
sns.lineplot(data=day_df, x="dteday", y="cnt", hue="year_name", ax=ax[0, 1])
ax[0, 1].set(title="Daily demand over time", xlabel="date", ylabel="cnt")
sns.barplot(data=day_df, x="season_name", y="cnt", order=["winter", "spring", "summer", "autumn"], ax=ax[1, 0])
ax[1, 0].set(title="Mean demand by season", xlabel="")
sns.boxplot(data=day_df, x="weather_name", y="cnt", order=["clear", "mist", "light rain/snow", "heavy rain/snow"], ax=ax[1, 1])
ax[1, 1].set(title="Demand by weather category", xlabel="")
fig.suptitle("Figure 2 - Daily demand, season and weather", y=1.01)
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_daily_eda.png", dpi=150); plt.show()

comp = day_df.groupby("season_name")[["casual", "registered"]].mean().reindex(["winter", "spring", "summer", "autumn"])
fig, ax = plt.subplots(1, 2, figsize=(13, 4.5))
comp.plot(kind="bar", stacked=True, ax=ax[0], color=["#ED553B", "#20639B"])
ax[0].set(title="Mean user-group composition by season", xlabel="", ylabel="mean daily rental count")
ax[0].tick_params(axis="x", rotation=0)
for c, color in zip(["temp", "atemp", "hum", "windspeed"], ["#20639B", "#3CAEA3", "#F6D55F", "#ED553B"]):
    ax[1].scatter(day_df[c], day_df.cnt, s=10, alpha=.28, label=c, color=color)
ax[1].set(title="Continuous weather measures and demand", xlabel="normalised weather value", ylabel="cnt")
ax[1].legend()
fig.suptitle("Figure 3 - User composition and continuous weather", y=1.02)
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_composition_weather.png", dpi=150); plt.show()

```

Figure 2 - Daily demand, season and weather





Figure 3 - User composition and continuous weather




Demand rises across the two years, peaks in warmer seasons and falls in adverse weather. Registered users provide most rentals, while the casual share expands in warmer seasons. Both user columns remain descriptive because they sum to the target.

A documentation correction. The dataset description states `season` (1:springer, 2:summer, 3:fall, 4:winter) . The dates disagree. Code 1 covers December to March at a mean of 12.2 degrees Celsius and the lowest mean demand of 2,604; code 3 covers June to September at 29.0 degrees and the highest mean demand of 5,644. The codes are therefore offset by one quarter from the published labels, and this notebook labels them `1 = winter, 2 = spring, 3 = summer, 4 = autumn` on the evidence of the calendar and temperature rather than the documentation. Nothing in the modelling changes, because `season` enters every pipeline as a categorical code and never as a name; the correction affects the readability of Figures 2, 3 and 13 and the wording of this section.

```
In [6]: weekday_lbl = {0: "Sun", 1: "Mon", 2: "Tue", 3: "Wed", 4: "Thu", 5: "Fri", 6: "Sat"}
        day_df["weekday_name"] = day_df.weekday.map(weekday_lbl)

        fig, ax = plt.subplots(1, 2, figsize=(13, 4.5))
        sns.lineplot(data=day_df, x="mnth", y="cnt", hue="year_name", marker="o", ax=ax[0])
        ax[0].set(title="Figure 4 - Mean rentals by month and year", xlabel="month", ylabel="mean cnt")
        ax[0].set_xticks(range(1, 13))
        sns.barplot(data=day_df, x="weekday_name", y="cnt",
                    order=["Sun", "Mon", "Tue", "Wed", "Thu", "Fri", "Sat"], ax=ax[1])
        ax[1].set(title="Figure 5 - Mean rentals by weekday", xlabel="weekday", ylabel="mean cnt")
        plt.tight_layout(); plt.savefig(FIG_DIR / "v4_month_weekday.png", dpi=160); plt.show()
```



Monthly demand follows the seasonal wave and every 2012 month exceeds its 2011 counterpart. Weekday means vary much less, showing that aggregation blends working-day and leisure patterns visible in the hourly profiles.

```
In [7]: corr_cols = ["season", "yr", "mnth", "holiday", "weekday", "workingday",
                    "weathersit", "temp", "atemp", "hum", "windspeed", "cnt"]
corr_matrix = day_df[corr_cols].corr()
cnt_correlations = corr_matrix["cnt"].drop("cnt").sort_values(ascending=False)
assert cnt_correlations.index[0] == "atemp"
assert cnt_correlations.loc["weathersit"] < 0 and cnt_correlations.loc["windspeed"] < 0

plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap="coolwarm", center=0,
            square=True, cbar_kws={"shrink": .8}, annot_kws={"size": 7})
plt.title("Figure 6 - Correlation matrix for daily variables")
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_correlation_heatmap.png", dpi=160); plt.show
print("Table 5 - Correlations with cnt")
display(cnt_correlations.round(3).to_frame("correlation_with_cnt"))
```





Table 5 - Correlations with cnt

	correlation_with_cnt
atemp	0.631
temp	0.627
yr	0.567
season	0.406
mnth	0.280
weekday	0.067
workingday	0.061
holiday	-0.068
hum	-0.115
windspeed	-0.235
weathersit	-0.297

`atemp` has the highest correlation with `cnt` (+0.631), followed by `temp` (+0.627) and `yr` (+0.567). Adverse weather (`weathersit` , -0.297), windspeed (-0.235) and corrected humidity (-0.115) are negative. The near-collinearity of `temp` and `atemp` cautions against interpreting linear coefficients independently; categorical codes also limit causal interpretation.

```
In [8]: # Stage 2.6: use seaborn's pairplot function and include its output.
pairplot_data = day_df[["temp", "atemp", "hum", "windspeed", "cnt"].dropna()
pair_grid = sns.pairplot(pairplot_data, corner=True, diag_kind="kde",
                        plot_kws={"alpha": .3, "s": 12})
pair_grid.fig.suptitle("Figure 7 - Seaborn pairplot of weather variables and daily demand")
pair_grid.savefig(FIG_DIR / "v4_pairplot.png", dpi=140, bbox_inches="tight")
plt.show()
```

Figure 7 - Seaborn pairplot of weather variables and daily demand



The pairplot confirms positive, nonlinear relationships between demand and both temperature measures. Humidity and windspeed show weaker negative patterns and broad scatter. `temp` and `atemp` move almost identically, consistent with their high correlation, while no single weather variable fully explains demand.

3. Data Preparation

Two frames prevent the temporal question from shrinking the primary sample.

`primary_df` contains all 731 days and no past-demand inputs. `temporal_df` starts after a causal seven-day warm-up and contains 724 days.

```
In [9]: # Shared calendar/weather engineering. Interactions use only target-day exogenous inputs.
day_df["atemp_hum"] = day_df["atemp"] * day_df["hum"]
day_df["temp_sq"] = day_df["temp"] ** 2
day_df["days_since_start"] = (day_df.dteday - day_df.dteday.min()).dt.days
day_df["mnth_sin"], day_df["mnth_cos"] = np.sin(2*np.pi*day_df.mnth/12), np.cos(2*np.pi*day_df.mnth/12)
day_df["weekday_sin"], day_df["weekday_cos"] = np.sin(2*np.pi*day_df.weekday/7), np.cos(2*np.pi*day_df.weekday/7)
day_df["season_sin"], day_df["season_cos"] = np.sin(2*np.pi*day_df.season/4), np.cos(2*np.pi*day_df.season/4)

primary_df = day_df.copy().reset_index(drop=True)

# Causal temporal features: every window ends before the target day.
day_df["lag_1_cnt"] = day_df.cnt.shift(1)
day_df["lag_7_cnt"] = day_df.cnt.shift(7)
day_df["roll_7_cnt"] = day_df.cnt.shift(1).rolling(7).mean()
assert day_df.lag_1_cnt.equals(day_df.cnt.shift(1))
assert day_df.lag_7_cnt.equals(day_df.cnt.shift(7))
for i in range(7, len(day_df)):
    assert np.isclose(day_df.loc[i, "roll_7_cnt"], day_df.loc[i-7:i-1, "cnt"].mean())
temporal_df = day_df.dropna(subset=["roll_7_cnt"]).reset_index(drop=True)

assert len(primary_df) == 731
assert len(temporal_df) == 724
print("primary_df rows:", len(primary_df), "| temporal_df rows:", len(temporal_df))
```

primary_df rows: 731 | temporal_df rows: 724

```

In [10]: PRIMARY_NOMINAL = ["season", "mnth", "weekday", "weathersit"]
PRIMARY_BINARY = ["yr", "holiday", "workingday"]
WEATHER = ["temp", "atemp", "hum", "windspeed"]
CYCLICAL = ["mnth_sin", "mnth_cos", "weekday_sin", "weekday_cos", "season_sin", "season_cos"]
INTERACTIONS = ["atemp_hum", "temp_sq"]
TREND = ["days_since_start"]
AUTOREGRESSIVE = ["lag_1_cnt", "lag_7_cnt", "roll_7_cnt"]

PRIMARY_FEATURE_SETS = {
    "core": PRIMARY_NOMINAL + PRIMARY_BINARY + WEATHER,
    "core + interactions": PRIMARY_NOMINAL + PRIMARY_BINARY + WEATHER + INTERACTIONS,
    "core + interactions + trend": PRIMARY_NOMINAL + PRIMARY_BINARY + WEATHER + INTERACTIONS + TREND
}
TEMPORAL_FEATURE_SETS = {
    "core": ["weathersit", "holiday", "workingday"] + WEATHER + CYCLICAL,
    "core + interactions": ["weathersit", "holiday", "workingday"] + WEATHER + CYCLICAL + INTERACTIONS,
    "core + interactions + autoregressive": ["weathersit", "holiday", "workingday"] + WEATHER + CYCLICAL + INTERACTIONS + AUTOREGRESSIVE,
    "core + interactions + autoregressive + trend": ["weathersit", "holiday", "workingday"] + WEATHER + CYCLICAL + INTERACTIONS + AUTOREGRESSIVE + TREND
}

for feature_list in list(PRIMARY_FEATURE_SETS.values()) + list(TEMPORAL_FEATURE_SETS.values()):
    assert "casual" not in feature_list and "registered" not in feature_list
    assert all(c not in cols for cols in PRIMARY_FEATURE_SETS.values() for c in AUTOREGRESSIVE)
    assert all("yr" not in cols for cols in TEMPORAL_FEATURE_SETS.values())

def make_preprocessor(columns, temporal=False):
    nominal = ["weathersit"] if temporal else [c for c in PRIMARY_NOMINAL if c in columns]
    binary_pool = ["holiday", "workingday"] + ([ "yr" ] if not temporal else [])
    binary = [c for c in binary_pool if c in columns]
    continuous = [c for c in columns if c not in nominal + binary]
    return ColumnTransformer([
        ("nominal", OneHotEncoder(drop="first", handle_unknown="ignore"), nominal),
        ("continuous", Pipeline([("impute", SimpleImputer(strategy="median")),
                                ("scale", StandardScaler())]), continuous),
        ("binary", "passthrough", binary),
    ])

def pipeline_for(model, columns, temporal=False):
    return Pipeline([("prep", make_preprocessor(columns, temporal=temporal)), ("model", model)])

def metrics(y_true, predictions):
    mse = mean_squared_error(y_true, predictions)
    return {"MAE": mean_absolute_error(y_true, predictions), "MSE": mse,
            "RMSE": float(np.sqrt(mse)), "R2": r2_score(y_true, predictions)}

GRIDS = {
    "Linear Regression": (LinearRegression(), {}),
    "K-Nearest Neighbors": (KNeighborsRegressor(), {
        "model__n_neighbors": [3, 5, 7, 9, 11, 15], "model__weights": ["uniform", "distance"]
    }),
    "Random Forest": (RandomForestRegressor(random_state=RANDOM_STATE, n_jobs=1), {
        "model__n_estimators": [200, 400], "model__max_depth": [None, 8, 16],
        "model__min_samples_leaf": [1, 2, 4]}),
    "Gradient Boosting": (GradientBoostingRegressor(random_state=RANDOM_STATE), {
        "model__n_estimators": [200, 400], "model__max_depth": [2, 3],
        "model__learning_rate": [0.05, 0.1]}),
}

MODEL_SIMPLICITY = {"Linear Regression": 1, "K-Nearest Neighbors": 2,
                    "Random Forest": 3, "Gradient Boosting": 4}
FEATURE_SIMPLICITY = {name: i for i, name in enumerate(
    ["core", "core + interactions", "core + interactions + trend",
    "core + interactions + autoregressive", "core + interactions + autoregressive + trend"]
)}

def compact_params(params):
    if not params:
        return "default"

```



```

labels = {"model__learning_rate": "lr", "model__max_depth": "depth",
          "model__n_estimators": "estimators", "model__min_samples_leaf": "leaf",
          "model__n_neighbors": "neighbors", "model__weights": "weights"}
return "; ".join(f"{labels.get(key, key)}={value}" for key, value in params.items())

def choose_configuration(rows):
    frame = pd.DataFrame(rows).copy()
    best_mae = frame["cv_mean_MAE"].min()
    shortlist = frame[frame["cv_mean_MAE"] <= 1.05 * best_mae].copy()
    shortlist["simplicity"] = shortlist["model"].map(MODEL_SIMPLICITY) + shortlist["feature"]
    return shortlist.sort_values(["cv_mean_RMSE", "simplicity", "cv_mean_MAE"]).iloc[0]

```

4. Modelling

Four regression families are compared with lightweight grids in scikit-learn (Pedregosa et al., 2011). Linear Regression tests a global additive relationship; K-Nearest Neighbors tests local similarity; Random Forest averages decorrelated trees; Gradient Boosting sequentially corrects residuals (Friedman, 2001). Scaling and fold-fitted median imputation live inside every pipeline.

Both experiments run over the same 731 days and differ in how those days are split and which inputs are permitted. Each produces one learned estimate and one reference estimate, giving the four columns compared in Table 13.

```
In [11]: from IPython.display import Image, display as show_image

# Figure 8 is authored in Mermaid and rendered to PNG so that it survives the notebook,
# HTML and PDF exports identically. The source is kept in the repository alongside it.
show_image(Image(filename=str(FIG_DIR / "v4_pipeline.png"), width=760))
print("Figure 8 - How the two experiments produce the four estimates compared in Table 13")
```




Figure 8 - How the two experiments produce the four estimates compared in Table 13

4.1 Primary model comparison: random 75/25 holdout

The split is created once with `random_state=42`. Only the 75% training partition participates in family, feature-set and hyperparameter selection. Each of the 12 family/feature-set combinations uses shuffled five-fold `KFold` and MAE scoring. The 25% holdout is accessed only after the global configuration and one configuration per family have been frozen.

```

In [12]: all_indices = primary_df.index.to_numpy()
primary_train_idx, primary_holdout_idx = train_test_split(
    all_indices, test_size=0.25, random_state=RANDOM_STATE
)
assert set(primary_train_idx).isdisjoint(primary_holdout_idx)
assert len(primary_train_idx) == 548 and len(primary_holdout_idx) == 183

primary_train = primary_df.loc[primary_train_idx].copy()
primary_holdout = primary_df.loc[primary_holdout_idx].copy()
primary_cv = KFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)
primary_rows, primary_estimators = [], {}

for feature_name, columns in PRIMARY_FEATURE_SETS.items():
    for model_name, (estimator, grid) in GRIDS.items():
        search = GridSearchCV(
            pipeline_for(clone(estimator), columns, temporal=False), grid,
            scoring="neg_mean_absolute_error", cv=primary_cv, n_jobs=1, refit=True
        )
        search.fit(primary_train[columns], primary_train.cnt)
        rmse_cv = cross_validate(
            search.best_estimator_, primary_train[columns], primary_train.cnt,
            cv=primary_cv, scoring="neg_root_mean_squared_error", n_jobs=1
        )["test_score"]
        row = {
            "model": model_name, "feature_set": feature_name,
            "cv_mean_MAE": -search.best_score_,
            "cv_std_MAE": search.cv_results_["std_test_score"][search.best_index_],
            "cv_mean_RMSE": -rmse_cv.mean(), "cv_std_RMSE": rmse_cv.std(),
            "best_params": compact_params(search.best_params_),
        }
        primary_rows.append(row)
        primary_estimators[(model_name, feature_name)] = search.best_estimator_

primary_cv_table = (pd.DataFrame(primary_rows)
                    .sort_values(["cv_mean_MAE", "cv_mean_RMSE"])
                    .reset_index(drop=True))
primary_cv_table.to_csv(OUTPUT_DIR / "primary_cv_selection_v4.csv", index=False)

primary_choice = choose_configuration(primary_rows)
primary_family_choices = {
    family: choose_configuration([r for r in primary_rows if r["model"] == family])
    for family in GRIDS
}
primary_winner_name = primary_choice["model"]
primary_winner_features = primary_choice["feature_set"]
primary_winner_columns = PRIMARY_FEATURE_SETS[primary_winner_features]
primary_winner = primary_estimators[(primary_winner_name, primary_winner_features)]

# Leakage audit: choices contain CV statistics from training indices only; the holdout has
# been passed to fit, GridSearchCV, cross_validate or choose_configuration.
selection_indices = set(primary_train_idx)
assert selection_indices.isdisjoint(set(primary_holdout_idx))
assert len(primary_df) == 731
print("Table 6 – Primary CV selection (training partition only)")
display(primary_cv_table.round({"cv_mean_MAE": 1, "cv_std_MAE": 1,
                                "cv_mean_RMSE": 1, "cv_std_RMSE": 1}))
print("Frozen primary configuration:", primary_winner_name, "|", primary_winner_features,
      "|", primary_choice["best_params"])

```

Table 6 – Primary CV selection (training partition only)

	model	feature_set	cv_mean_MAE	cv_std_MAE	cv_mean_RMSE	cv_std_RMSE	best_para
0	Gradient Boosting	core + interactions + trend	477.0	34.1	657.6	45.4	lr=0.1; depth estimators=4
1	Random Forest	core + interactions + trend	498.0	18.1	723.5	38.5	depth=16; leaf estimators=2
2	Gradient Boosting	core	513.8	56.4	690.8	53.2	lr=0.05; depth estimators=4
3	Gradient Boosting	core + interactions	519.5	47.4	697.1	46.4	lr=0.05; depth estimators=4
4	Random Forest	core	533.6	26.5	746.8	25.7	depth=16; leaf estimators=4
5	Random Forest	core + interactions	539.1	24.3	752.6	23.2	depth=16; leaf estimators=4
6	Linear Regression	core + interactions + trend	559.3	19.8	752.4	36.9	default
7	Linear Regression	core + interactions	560.4	19.3	752.6	38.5	default
8	K-Nearest Neighbors	core + interactions + trend	563.0	15.1	794.0	40.0	neighbors: weights=distance
9	Linear Regression	core	591.0	14.7	806.8	41.9	default
10	K-Nearest Neighbors	core + interactions	756.5	39.1	966.5	48.0	neighbors: weights=distance
11	K-Nearest Neighbors	core	760.7	42.7	957.3	52.9	neighbors: weights=distance

Frozen primary configuration: Gradient Boosting | core + interactions + trend | lr=0.1; depth=2; estimators=400

```

In [13]: # First use of the fixed 25% holdout: score one training-selected configuration per family
primary_holdout_rows = []
mean_model = DummyRegressor(strategy="mean").fit(
    np.zeros((len(primary_train), 1)), primary_train.cnt
)
baseline_pred = mean_model.predict(np.zeros((len(primary_holdout), 1)))
primary_holdout_rows.append({"model": "Mean baseline", "feature_set": "constant training r",
                             "best_params": "default", **metrics(primary_holdout.cnt, base

primary_family_predictions = {}
for family, choice in primary_family_choices.items():
    feature_name = choice["feature_set"]
    columns = PRIMARY_FEATURE_SETS[feature_name]
    fitted = primary_estimators[(family, feature_name)]
    pred = fitted.predict(primary_holdout[columns])
    primary_family_predictions[family] = pred
    primary_holdout_rows.append({"model": family, "feature_set": feature_name,
                                 "best_params": choice["best_params"],
                                 **metrics(primary_holdout.cnt, pred)})

primary_holdout_table = (pd.DataFrame(primary_holdout_rows)
                          .sort_values("MAE").reset_index(drop=True))
primary_holdout_table.to_csv(OUTPUT_DIR / "primary_holdout_metrics_v4.csv", index=False)
primary_metrics = primary_holdout_table.loc[
    primary_holdout_table.model.eq(primary_winner_name)
].iloc[0]
primary_baseline = primary_holdout_table.loc[
    primary_holdout_table.model.eq("Mean baseline")
].iloc[0]
primary_improvement = 1 - primary_metrics.MAE / primary_baseline.MAE
primary_mae_pct = primary_metrics.MAE / primary_holdout.cnt.mean()
primary_summary = pd.DataFrame([
    "selected_model": primary_winner_name,
    "feature_set": primary_winner_features,
    "best_params": primary_choice["best_params"],
    "holdout_MAE": primary_metrics.MAE,
    "holdout_RMSE": primary_metrics.RMSE,
    "holdout_R2": primary_metrics.R2,
    "mean_baseline_MAE": primary_baseline.MAE,
    "baseline_improvement_pct": 100 * primary_improvement,
    "MAE_pct_of_holdout_mean_demand": 100 * primary_mae_pct,
])
primary_summary.to_csv(OUTPUT_DIR / "primary_summary_v4.csv", index=False)
print("Table 7 - Frozen family configurations on the primary holdout")
display(primary_holdout_table.round({"MAE": 1, "MSE": 1, "RMSE": 1, "R2": 3}))
print("Table 8 - Final selected-model summary")
display(primary_summary.round(3))

```

Table 7 - Frozen family configurations on the primary holdout

	model	feature_set	best_params	MAE	MSE	RMSE	R2
0	Gradient Boosting	core + interactions + trend	lr=0.1; depth=2; estimators=400	433.9	397608.0	630.6	0.897
1	Random Forest	core + interactions + trend	depth=16; leaf=1; estimators=200	446.6	463735.1	681.0	0.880
2	K-Nearest Neighbors	core + interactions + trend	neighbors=7; weights=distance	501.0	612073.9	782.4	0.841
3	Linear Regression	core + interactions + trend	default	529.9	524664.9	724.3	0.864
4	Mean baseline	constant training mean	default	1663.0	3933577.2	1983.3	-0.021

Table 8 - Final selected-model summary

	selected_model	feature_set	best_params	holdout_MAE	holdout_RMSE	holdout_R2	mean_base
0	Gradient Boosting	core + interactions + trend	lr=0.1; depth=2; estimators=400	433.863	630.562	0.897	

```
In [14]: # Sensitivity of the frozen primary winner to retaining the raw hum=0 record.
raw_primary = primary_df.copy()
raw_primary.loc[raw_primary.dteday.eq(zero_hum_date), "hum"] = 0.0
raw_primary["atemp_hum"] = raw_primary.atemp * raw_primary.hum
sens_model = clone(primary_winner).fit(
    raw_primary.loc[primary_train_idx, primary_winner_columns],
    raw_primary.loc[primary_train_idx, "cnt"]
)
sens_pred = sens_model.predict(raw_primary.loc[primary_holdout_idx, primary_winner_columns])
corrected_pred = primary_winner.predict(primary_holdout[primary_winner_columns])
humidity_sensitivity = pd.DataFrame([
    {"treatment": "hum=0 marked missing; fold-fitted median imputation",
     **metrics(primary_holdout.cnt, corrected_pred)},
    {"treatment": "raw hum=0 retained", **metrics(primary_holdout.cnt, sens_pred)},
])
humidity_sensitivity.to_csv(OUTPUT_DIR / "humidity_sensitivity_v4.csv", index=False)
print("Table 9 - Humidity correction sensitivity on the frozen primary model")
display(humidity_sensitivity.round({"MAE": 2, "MSE": 1, "RMSE": 2, "R2": 4}))
```

Table 9 – Humidity correction sensitivity on the frozen primary model

	treatment	MAE	MSE	RMSE	R2
0	hum=0 marked missing; fold-fitted median imput...	433.86	397608.0	630.56	0.8968
1	raw hum=0 retained	430.85	397549.6	630.52	0.8968

Primary result. Gradient Boosting with interactions and elapsed-time trend was frozen at CV MAE 477.0 (SD 34.1) and RMSE 657.6. Random Forest entered the 5% shortlist at MAE 498.0 but had higher RMSE (723.5). The selected grid point used learning rate 0.1, depth 2 and 400 estimators.

On the untouched 183-day holdout it reached MAE 433.9, RMSE 630.6 and R-squared 0.897. MAE was 73.9% below the baseline's 1,663.0 and equalled **10.1% of mean holdout demand**, meeting the 40% criterion. Random Forest reached 446.6, KNN 501.0 and Linear Regression 529.9; the winner remained frozen.

Marking zero humidity missing produced MAE 433.9 versus 430.9 when retained. The sub-1% difference leaves the conclusion unchanged; fold-fitted imputation preserves the physically plausible treatment.

4.2 Secondary temporal robustness check

Selection now uses `TimeSeriesSplit(n_splits=3)` inside 2011. The temporal design excludes `yr`, which is constant in the training year. It compares core cyclical calendar/weather inputs, interactions, causal recent-demand inputs and an optional elapsed-time trend. One configuration per family and the global temporal configuration are frozen before any 2012 score is produced.

```

In [15]: temporal_train = temporal_df[temporal_df.yr.eq(0)].copy().reset_index(drop=True)
temporal_test = temporal_df[temporal_df.yr.eq(1)].copy().reset_index(drop=True)
assert temporal_train.dteday.max() < pd.Timestamp("2012-01-01")
assert temporal_test.dteday.min() >= pd.Timestamp("2012-01-01")
assert all("yr" not in columns for columns in TEMPORAL_FEATURE_SETS.values())

temporal_cv = TimeSeriesSplit(n_splits=3)
fold_table = []
for fold, (tr, va) in enumerate(temporal_cv.split(temporal_train), 1):
    fold_table.append({
        "fold": fold, "train_rows": len(tr), "validation_rows": len(va),
        "train_window": f"{temporal_train.dteday.iloc[tr[0]].date()} to {temporal_train.dteday.iloc[tr[-1]].date()}",
        "validation_window": f"{temporal_train.dteday.iloc[va[0]].date()} to {temporal_train.dteday.iloc[va[-1]].date()}"
    })
print("Table 10 - TimeSeriesSplit folds within 2011")
display(pd.DataFrame(fold_table))

temporal_rows, temporal_estimators = [], {}
for feature_name, columns in TEMPORAL_FEATURE_SETS.items():
    for model_name, (estimator, grid) in GRIDS.items():
        search = GridSearchCV(
            pipeline_for(clone(estimator), columns, temporal=True), grid,
            scoring="neg_mean_absolute_error", cv=temporal_cv, n_jobs=1, refit=True
        )
        search.fit(temporal_train[columns], temporal_train.cnt)
        rmse_cv = cross_validate(
            search.best_estimator_, temporal_train[columns], temporal_train.cnt,
            cv=temporal_cv, scoring="neg_root_mean_squared_error", n_jobs=1
        )["test_score"]
        row = {"model": model_name, "feature_set": feature_name,
            "cv_mean_MAE": -search.best_score_,
            "cv_std_MAE": search.cv_results_["std_test_score"][search.best_index_],
            "cv_mean_RMSE": -rmse_cv.mean(), "cv_std_RMSE": rmse_cv.std(),
            "best_params": compact_params(search.best_params_)}
        temporal_rows.append(row)
        temporal_estimators[(model_name, feature_name)] = search.best_estimator_

temporal_cv_table = (pd.DataFrame(temporal_rows)
    .sort_values(["cv_mean_MAE", "cv_mean_RMSE"]).reset_index(drop=True))
temporal_cv_table.to_csv(OUTPUT_DIR / "temporal_cv_selection_v4.csv", index=False)
temporal_choice = choose_configuration(temporal_rows)
temporal_family_choices = {
    family: choose_configuration([r for r in temporal_rows if r["model"] == family])
    for family in GRIDS
}
temporal_winner_name = temporal_choice["model"]
temporal_winner_features = temporal_choice["feature_set"]
temporal_winner = temporal_estimators[(temporal_winner_name, temporal_winner_features)]

# Selection records and fitted estimators contain 2011 only. 2012 has not been scored yet.
assert temporal_train.yr.eq(0).all()
assert not temporal_train.dteday.isin(temporal_test.dteday).any()
print("Table 11 - Temporal CV selection (2011 only)")
display(temporal_cv_table.round({"cv_mean_MAE": 1, "cv_std_MAE": 1,
    "cv_mean_RMSE": 1, "cv_std_RMSE": 1}))
print("Frozen temporal configuration:", temporal_winner_name, "|", temporal_winner_features,
    "|", temporal_choice["best_params"])

```

Table 10 - TimeSeriesSplit folds within 2011

	fold	train_rows	validation_rows	train_window	validation_window
0	1	91	89	2011-01-08 to 2011-04-08	2011-04-09 to 2011-07-06
1	2	180	89	2011-01-08 to 2011-07-06	2011-07-07 to 2011-10-03
2	3	269	89	2011-01-08 to 2011-10-03	2011-10-04 to 2011-12-31

Table 11 – Temporal CV selection (2011 only)

	model	feature_set	cv_mean_MAE	cv_std_MAE	cv_mean_RMSE	cv_std_RMSE	best_parameters
0	Linear Regression	core + interactions + autoregressive	602.5	68.5	763.2	82.4	default
1	Linear Regression	core + interactions + autoregressive + trend	665.8	49.4	815.2	50.8	default
2	Linear Regression	core + interactions	763.9	128.1	929.9	96.0	default
3	Gradient Boosting	core	878.0	356.6	1008.2	341.8	lr=0.1; depth=1; k=1; neighbors=1; weights=dis
4	Gradient Boosting	core + interactions	889.5	371.3	1025.5	351.1	lr=0.1; depth=1; k=1; neighbors=1; weights=dis
5	Gradient Boosting	core + interactions + autoregressive + trend	900.0	459.1	1086.4	406.9	lr=0.05; depth=1; k=1; neighbors=1; weights=dis
6	Gradient Boosting	core + interactions + autoregressive	905.0	459.5	1059.6	425.8	lr=0.1; depth=1; k=1; neighbors=1; weights=dis
7	Linear Regression	core	917.3	321.5	1072.1	299.1	default
8	Random Forest	core + interactions + autoregressive + trend	922.1	410.7	1094.3	368.3	depth=1; k=1; neighbors=1; weights=dis
9	Random Forest	core + interactions + autoregressive	924.4	401.3	1089.4	357.7	depth=1; k=1; neighbors=1; weights=dis
10	Random Forest	core	962.3	386.2	1111.3	355.9	depth=8; k=1; neighbors=1; weights=dis
11	Random Forest	core + interactions	974.3	381.9	1121.1	352.8	depth=1; k=1; neighbors=1; weights=dis
12	K-Nearest Neighbors	core	1056.4	349.3	1269.4	281.8	neighbors=1; weights=dis
13	K-Nearest Neighbors	core + interactions	1068.6	335.8	1269.9	266.2	neighbors=1; weights=dis
14	K-Nearest Neighbors	core + interactions + autoregressive + trend	1093.2	553.9	1314.7	483.4	neighbors=1; weights=dis
15	K-Nearest Neighbors	core + interactions + autoregressive	1123.7	535.1	1327.1	476.2	neighbors=1; weights=dis

Frozen temporal configuration: Linear Regression | core + interactions + autoregressive | default

```

In [16]: # First 2012 scoring: frozen family candidates and causal baselines use identical dates.
temporal_rows_2012, temporal_predictions = [], {}
for family, choice in temporal_family_choices.items():
    feature_name = choice["feature_set"]
    columns = TEMPORAL_FEATURE_SETS[feature_name]
    fitted = temporal_estimators[(family, feature_name)]
    pred = fitted.predict(temporal_test[columns])
    temporal_predictions[family] = pred
    temporal_rows_2012.append({"model": family, "feature_set": feature_name,
                              "best_params": choice["best_params"],
                              **metrics(temporal_test.cnt, pred)})

temporal_baseline_predictions = {
    "Naive lag-1": temporal_test.lag_1_cnt.to_numpy(),
    "Naive lag-7": temporal_test.lag_7_cnt.to_numpy(),
    "Naive rolling-7": temporal_test.roll_7_cnt.to_numpy(),
    "2011 constant mean": np.full(len(temporal_test), temporal_train.cnt.mean()),
}
for name, pred in temporal_baseline_predictions.items():
    temporal_predictions[name] = pred
    temporal_rows_2012.append({"model": name, "feature_set": "causal baseline",
                              "best_params": "default", **metrics(temporal_test.cnt, pred)})

temporal_holdout_table = (pd.DataFrame(temporal_rows_2012)
                          .sort_values("MAE").reset_index(drop=True))
temporal_holdout_table.to_csv(OUTPUT_DIR / "temporal_holdout_metrics_v4.csv", index=False)
temporal_selected_metrics = temporal_holdout_table.loc[
    temporal_holdout_table.model.eq(temporal_winner_name)
].iloc[0]
rolling_metrics = temporal_holdout_table.loc[
    temporal_holdout_table.model.eq("Naive rolling-7")
].iloc[0]
temporal_advantage = 1 - temporal_selected_metrics.MAE / rolling_metrics.MAE
print("Table 12 – Frozen temporal candidates and baselines on 2012")
display(temporal_holdout_table.round({"MAE": 1, "MSE": 1, "RMSE": 1, "R2": 3}))

```

Table 12 – Frozen temporal candidates and baselines on 2012

	model	feature_set	best_params	MAE	MSE	RMSE	R2
0	Naive rolling-7	causal baseline	default	847.8	1388485.8	1178.3	0.565
1	Naive lag-1	causal baseline	default	870.2	1553433.8	1246.4	0.513
2	Linear Regression	core + interactions + autoregressive	default	1047.4	1449639.0	1204.0	0.546
3	Naive lag-7	causal baseline	default	1110.4	2439375.1	1561.8	0.235
4	Random Forest	core + interactions + autoregressive	depth=None; leaf=1; estimators=400	1673.3	3462504.7	1860.8	-0.085
5	Gradient Boosting	core	lr=0.1; depth=2; estimators=400	2040.5	4834520.0	2198.8	-0.515
6	K-Nearest Neighbors	core	neighbors=3; weights=distance	2077.1	5132081.4	2265.4	-0.609
7	2011 constant mean	causal baseline	default	2460.2	7829760.2	2798.2	-1.454

Temporal result. Linear Regression with interactions and causal autoregressive inputs was frozen from 2011 at CV MAE 602.5 and CV RMSE 763.2. Adding elapsed-time trend raised CV MAE to 665.8, so the trend was excluded before 2012 evaluation. This contrast with the primary result shows that elapsed time helps interpolate within a mixed-year domain but does not establish a transferable growth law.

In 2012, the frozen Linear Regression reached MAE 1,047.4, RMSE 1,204.0 and R-squared 0.546. Rolling-7 reached MAE 847.8, RMSE 1,178.3 and R-squared 0.565, leaving the selected model 23.5% worse on MAE. Random Forest scored MAE 1,673.3, while the 2011-selected Gradient Boosting and KNN configurations exceeded 2,000. Forward temporal robustness therefore fails.

A rolling one-day-ahead test can update lags after every observed day. A forecast issued for all of 2012 on 1 January could not use those later counts and would be a different experiment with a longer horizon and recursively unavailable inputs.

4.3 Worked example

Aggregate error statistics hide what a prediction actually is. Table 13 takes six holdout dates that both experiments scored, and shows the same day estimated four ways: by the frozen primary model, by the constant-mean reference it must beat, by the frozen temporal model, and by the rolling seven-day rule it must beat. Reading one row left to right is the shortest description of this entire report.

```

In [17]: # Table 13 – the same dates scored by every approach compared in this report.
primary_view = primary_holdout.assign(
    primary_prediction=primary_winner.predict(primary_holdout[primary_winner_columns]),
    mean_reference=primary_train.cnt.mean(),
)
temporal_view = temporal_test[["dteday"]].assign(
    temporal_prediction=temporal_predictions[temporal_winner_name],
    rolling7_reference=temporal_predictions["Naive rolling-7"],
)
shared = primary_view.merge(temporal_view, on="dteday", how="inner").sort_values("dteday")

# Six evenly spaced dates across the shared period, so the sample is reproducible.
positions = np.linspace(0, len(shared) - 1, 6).round().astype(int)
sample = shared.iloc[positions]

season_lbl = {1: "winter", 2: "spring", 3: "summer", 4: "autumn"}
weather_lbl = {1: "clear", 2: "mist", 3: "light rain/snow", 4: "severe"}
worked = pd.DataFrame({
    "date": sample.dteday.dt.strftime("%Y-%m-%d"),
    "season": sample.season.map(season_lbl),
    "day": sample.dteday.dt.strftime("%a"),
    "weather": sample.weathersit.map(weather_lbl),
    "temp_C": (sample.temp * 41).round(1),
    "humidity_pct": (sample.hum * 100).round(0),
    "ACTUAL": sample.cnt.astype(int),
    f"{primary_winner_name} (primary)": sample.primary_prediction.round(0).astype(int),
    "mean reference": sample.mean_reference.round(0).astype(int),
    f"{temporal_winner_name} (temporal)": sample.temporal_prediction.round(0).astype(int),
    "rolling-7 reference": sample.rolling7_reference.round(0).astype(int),
}).reset_index(drop=True)
worked.to_csv(OUTPUT_DIR / "worked_example_v4.csv", index=False)

print("Table 13 – Six shared dates scored by every approach")
print("Each row is one day: inputs on the left, the realised count, then four estimates of")
display(worked)

for label, column in [(f"{primary_winner_name} (primary)", "primary_prediction"),
                      ("mean reference", "mean_reference"),
                      (f"{temporal_winner_name} (temporal)", "temporal_prediction"),
                      ("rolling-7 reference", "rolling7_reference")]:
    print(f" mean absolute error on these six dates, {label:34s}: "
          f"{np.abs(sample.cnt - sample[column]).mean():7.0f}")
print(f"\nShared dates available: {len(shared)} (2012 days that are also primary holdout)")

```

Table 13 – Six shared dates scored by every approach

Each row is one day: inputs on the left, the realised count, then four estimates of it.

	date	season	day	weather	temp_C	humidity_pct	ACTUAL	Gradient Boosting (primary)	mean reference	Linear Regression (temporal)
0	2012-01-04	winter	Wed	mist	4.4	41.0	2368	2405	4575	1131
1	2012-03-19	winter	Mon	clear	22.3	73.0	6153	6596	4575	5114
2	2012-06-07	spring	Thu	clear	24.7	57.0	7494	7373	4575	6142
3	2012-08-03	summer	Fri	mist	31.4	64.0	7175	6422	4575	5495
4	2012-10-12	autumn	Fri	clear	17.9	54.0	7282	7453	4575	5769
5	2012-12-30	winter	Sun	clear	10.5	48.0	1796	1557	4575	1579

mean absolute error on these six dates, Gradient Boosting (primary)	:	294
mean absolute error on these six dates, mean reference	:	2465
mean absolute error on these six dates, Linear Regression (temporal)	:	1173
mean absolute error on these six dates, rolling-7 reference	:	299

Shared dates available: 80 (2012 days that are also primary holdout days).

Six days cannot establish anything, and the criteria in Section 5 are computed across the full 183-day holdout and the full 366-day 2012 period for that reason. What the rows do show is the shape of each approach.

The primary model tracks the realised count in both directions, from a 2,368-rental January day to a 7,494-rental June day. The constant reference cannot move: it answers 4,575 to every question, which is why its error grows with distance from the annual mean. The temporal model follows the direction of change but sits below the realised count on the growth days, because it was fitted on a smaller 2011 system and has no mechanism for the expansion that followed. The rolling seven-day rule carries no model at all and stays close, because each new observed day re-anchors the next estimate.

The primary and rolling-7 columns must not be read as a head-to-head contest. They belong to different comparisons in Table 1: the primary model was fitted on days drawn from both years, including days adjacent to these, and answers a conditional question; the rolling rule answers a forward one and needs no training. Their apparent similarity on six shared dates is a property of this sample, not a finding. The comparison that matters for each is in Tables 7 and 12 respectively.

5. Evaluation

5.1 Criteria summary and paired temporal description

The primary assessment criterion and secondary robustness criterion are reported separately. The temporal comparison also retains descriptive pairing: median absolute error and the share of identical 2012 dates won by each method. These summaries do not require an independence assumption. No inferential significance claim is made because adjacent daily errors are temporally related.

```
In [18]: primary_pass = primary_improvement >= 0.40
temporal_pass = temporal_advantage >= 0.05
evaluation_summary = pd.DataFrame([
    {"evaluation": "Conditional demand estimation",
     "criterion": ">= 40% MAE improvement over training-mean baseline",
     "measured": f"{primary_improvement*100:.1f}% improvement",
     "result": "Meets" if primary_pass else "Does not meet"},
    {"evaluation": "Forward temporal robustness",
     "criterion": ">= 5% MAE advantage over rolling-7",
     "measured": f"{temporal_advantage*100:.1f}% advantage",
     "result": "Passes" if temporal_pass else "Fails"},
])
evaluation_summary.to_csv(OUTPUT_DIR / "evaluation_summary_v4.csv", index=False)
print("Table 14 - Evaluation summary")
display(evaluation_summary)

y2012 = temporal_test.cnt.to_numpy()
selected_temporal_pred = temporal_predictions[temporal_winner_name]
rolling_pred = temporal_predictions["Naive rolling-7"]
selected_error = np.abs(y2012 - selected_temporal_pred)
rolling_error = np.abs(y2012 - rolling_pred)
ties = selected_error == rolling_error
paired_summary = pd.DataFrame([
    {"series": temporal_winner_name, "median_absolute_error": np.median(selected_error),
     "days_closer": int((selected_error < rolling_error).sum()),
     "win_rate_excluding_ties_pct": 100 * (selected_error < rolling_error).sum() / (~ties).sum()},
    {"series": "Naive rolling-7", "median_absolute_error": np.median(rolling_error),
     "days_closer": int((rolling_error < selected_error).sum()),
     "win_rate_excluding_ties_pct": 100 * (rolling_error < selected_error).sum() / (~ties).sum()}
])
paired_summary["tied_dates"] = int(ties.sum())
paired_summary.to_csv(OUTPUT_DIR / "paired_temporal_comparison_v4.csv", index=False)
print("Table 15 - Paired descriptive comparison on identical 2012 dates")
display(paired_summary.round(1))
```

Table 14 - Evaluation summary

	evaluation	criterion	measured	result
0	Conditional demand estimation	>= 40% MAE improvement over training-mean baseline	73.9% improvement	Meets
1	Forward temporal robustness	>= 5% MAE advantage over rolling-7	-23.5% advantage	Fails

Table 15 - Paired descriptive comparison on identical 2012 dates



	series	median_absolute_error	days_closer	win_rate_excluding_ties_pct	tied_dates
0	Linear Regression	1045.2	116	31.7	0
1	Naive rolling-7	636.1	250	68.3	0

```
In [19]: train_ceiling = int(temporal_train.cnt.max())
ceiling_rows = []
for name in list(GRIDS) + ["Naive rolling-7"]:
    pred = temporal_predictions[name]
    ceiling_rows.append({
        "model": name,
        "largest_2012_prediction": float(np.max(pred)),
        "2011_training_ceiling": train_ceiling,
        "2012_actual_maximum": int(temporal_test.cnt.max()),
        "mean_residual_actual_minus_prediction": float(np.mean(y2012 - pred)),
    })
extrapolation_ceiling = pd.DataFrame(ceiling_rows)
extrapolation_ceiling.to_csv(OUTPUT_DIR / "extrapolation_ceiling_v4.csv", index=False)
print("Table 16 - Extrapolation ceiling analysis")
display(extrapolation_ceiling.round(1))

fig, ax = plt.subplots(1, 2, figsize=(14, 5))
ax[0].plot(temporal_test.dteday, y2012, lw=.9, color="#20639B", label="actual 2012")
ax[0].plot(temporal_test.dteday, selected_temporal_pred, lw=.9, color="#ED553B",
            label=f"frozen {temporal_winner_name}")
ax[0].plot(temporal_test.dteday, rolling_pred, lw=.9, color="#3CAEA3", label="rolling-7")
ax[0].axhline(train_ceiling, color="black", ls="--", lw=1, label=f"2011 ceiling = {train_c")
ax[0].set(title="Figure 9 - Rolling day-ahead predictions in 2012", xlabel="date", ylabel="cnt")
ax[0].legend(fontsize=8)
ax[1].scatter(selected_temporal_pred, y2012-selected_temporal_pred, s=14, alpha=.4, color="#ED553B")
ax[1].axhline(0, color="black", ls="--")
ax[1].set(title=f"Figure 10 - {temporal_winner_name} temporal residuals",
            xlabel="predicted cnt", ylabel="actual - predicted")
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_temporal_robustness.png", dpi=160); plt.show
```

Table 16 - Extrapolation ceiling analysis

	model	largest_2012_prediction	2011_training_ceiling	2012_actual_maximum	mean_residual_a
0	Linear Regression	8247.5	6043	8714	
1	K-Nearest Neighbors	5540.7	6043	8714	
2	Random Forest	5567.2	6043	8714	
3	Gradient Boosting	5490.7	6043	8714	
4	Naive rolling-7	7988.0	6043	8714	



Tree ensembles predict averages within terminal regions and therefore have an extrapolation ceiling tied to observed training targets. The 2011 maximum was 6,043 rentals, while 2012 reached 8,714. Random Forest's largest prediction was 5,567 and Gradient Boosting's was 5,491; their mean under-predictions were 1,568 and 1,995 rentals per day. KNN also stayed below the training ceiling.

Linear Regression extended to 8,248 and reduced mean under-prediction to 832, explaining why it led the 2011-selected ML candidates. Rolling-7 reached 7,988 and had a mean residual of -2.5 because each new observation re-anchored the next prediction. Conditional accuracy within a mixed-year sample therefore did not transfer to the later growth period.

5.2 Explaining the frozen primary model

Permutation importance is model-agnostic and measures the increase in holdout MAE after one raw input is shuffled. It is computed only after model selection and does not change the winner. Contributions describe predictive associations in this dataset, not causal effects. TreeSHAP (Lundberg & Lee, 2017) is added only when the frozen primary family is Random Forest or Gradient Boosting.


```
In [20]: perm = permutation_importance(
    primary_winner, primary_holdout[primary_winner_columns], primary_holdout.cnt,
    scoring="neg_mean_absolute_error", n_repeats=30, random_state=RANDOM_STATE, n_jobs=1
)
permutation_table = (pd.DataFrame({
    "feature": primary_winner_columns,
    "MAE_increase_mean": perm.importances_mean,
    "MAE_increase_std": perm.importances_std,
}).sort_values("MAE_increase_mean", ascending=False).reset_index(drop=True))
permutation_table.to_csv(OUTPUT_DIR / "permutation_importance_v4.csv", index=False)
print("Table 17 - Model-agnostic permutation importance on the primary holdout")
display(permutation_table.head(12).round(2))

top_perm = permutation_table.head(12).sort_values("MAE_increase_mean")
fig, ax = plt.subplots(figsize=(9, 5.5))
ax.barh(top_perm.feature, top_perm.MAE_increase_mean, xerr=top_perm.MAE_increase_std,
        color="#20639B", alpha=.85)
ax.set(title=f"Figure 11 - Permutation importance: {primary_winner_name}",
        xlabel="increase in holdout MAE after permutation", ylabel="")
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_permutation_importance.png", dpi=160); plt.c
```

Table 17 - Model-agnostic permutation importance on the primary holdout

	feature	MAE_increase_mean	MAE_increase_std
0	days_since_start	950.28	53.08
1	atemp	409.77	41.65
2	hum	144.47	18.57
3	temp_sq	110.50	16.63
4	mnth	81.71	15.54
5	windspeed	76.41	15.94
6	weathersit	69.47	9.73
7	workingday	34.08	9.05
8	atemp_hum	29.41	6.80
9	temp	24.92	8.24
10	season	24.35	5.03
11	weekday	23.60	7.73

Figure 11 - Permutation importance: Gradient Boosting




```

In [21]: shap_table = None
if primary_winner_name in {"Random Forest", "Gradient Boosting"}:
    import shap
    fitted_prep = primary_winner.named_steps["prep"]
    transformed = fitted_prep.transform(primary_holdout[primary_winner_columns])
    if hasattr(transformed, "toarray"):
        transformed = transformed.toarray()
    transformed_names = fitted_prep.get_feature_names_out()
    explainer = shap.TreeExplainer(primary_winner.named_steps["model"])
    shap_values = explainer.shap_values(transformed, check_additivity=True)
    shap_table = (pd.DataFrame({"feature": transformed_names,
                               "mean_absolute_SHAP": np.abs(shap_values).mean(axis=0)})
                  .sort_values("mean_absolute_SHAP", ascending=False).reset_index(drop=True))
    shap_table.to_csv(OUTPUT_DIR / "shap_global_importance_v4.csv", index=False)
    print("Table 18 - TreeSHAP global importance for the frozen primary model")
    display(shap_table.head(12).round(2))
    shap.summary_plot(shap_values, transformed, feature_names=transformed_names,
                      max_display=12, show=False)
    plt.title(f"Figure 12 - TreeSHAP summary: {primary_winner_name}")
    plt.tight_layout(); plt.savefig(FIG_DIR / "v4_shap_summary.png", dpi=160, bbox_inches=
else:
    print("TreeSHAP omitted: the frozen primary winner is", primary_winner_name)



```

Table 18 - TreeSHAP global importance for the frozen primary model

	feature	mean_absolute_SHAP
0	continuous__days_since_start	1018.06
1	continuous__atemp	586.20
2	continuous__hum	234.58
3	continuous__temp_sq	201.12
4	continuous__windspeed	136.82
5	continuous__temp	119.90
6	nominal__weathersit_2	97.63
7	binary__workingday	87.97
8	nominal__weekday_6	69.75
9	nominal__mnth_8	55.72
10	nominal__weathersit_3	55.56
11	continuous__atemp_hum	52.08



```
In [22]: primary_pred = primary_winner.predict(primary_holdout[primary_winner_columns])
fig, ax = plt.subplots(1, 2, figsize=(13, 5))
ax[0].scatter(primary_holdout.cnt, primary_pred, alpha=.45, s=18, color="#20639B")
lims = [min(primary_holdout.cnt.min(), primary_pred.min()),
        max(primary_holdout.cnt.max(), primary_pred.max())]
ax[0].plot(lims, lims, "k--")
ax[0].set(title="Figure 13 - Primary holdout: predicted vs actual",
          xlabel="actual cnt", ylabel="predicted cnt")
ax[1].scatter(primary_pred, primary_holdout.cnt - primary_pred, alpha=.45, s=18, color="#ED7D31")
ax[1].axhline(0, color="black", ls="--")
ax[1].set(title="Figure 14 - Primary holdout residuals",
          xlabel="predicted cnt", ylabel="actual - predicted")
plt.tight_layout(); plt.savefig(FIG_DIR / "v4_primary_diagnostics.png", dpi=160); plt.show()
```



5.3 Interpretation

Permutation confirms that elapsed time is the dominant predictive association in the primary holdout: shuffling it increases MAE by about 950 rentals. Feels-like temperature follows at 410, then humidity at 144 and squared temperature at 111. TreeSHAP produces the same leading order, with mean absolute contributions of about 1,018, 586, 235 and 201 rentals. These values explain predictions across the observed sample; they do not imply that advancing time causes demand or that weather effects are causal.

The primary result is MAE 433.9, RMSE 630.6 and R-squared 0.897: 73.9% better than the baseline and 10.1% of mean holdout demand. Temporally, Linear Regression lost to rolling-7 by 23.5%; median error was 1,045 versus 636, with wins on 116 versus 250 dates. This paired description makes no independence claim.

6. Lessons Learned

6.1 What went well

The method built during Assessments 1 and 2 carried straight into this project and shortened the setup considerably. Declaring criteria before modelling, keeping a holdout untouched until selection finishes, running an ablation to attribute a metric change to a cause, and operating the notebook as a submission artefact were all decided in earlier assessments, so this one reached a working pipeline much faster. The remaining effort went into the questions specific to this dataset instead of into process.

The design also held up under its own checks. All 731 daily units remain available for the primary question, the holdout is isolated from every selection decision, and four families compete across three exogenous feature sets. The hourly file validates every daily total and exposes both intraday peaks and panel imbalance.

6.2 Challenges

Understanding the data took longer than fitting the models. Two files describe the same system at different granularities, `casual` and `registered` sum to the target and cannot be predictors, and one humidity reading is physically impossible. Diagnosing that reading required cross-checking the hourly file, and the correction then had to preserve sample size and stay inside the pipeline folds.

The concept that took longest to internalise is the difference between interpolation and extrapolation. A tree fits both years comfortably when they are shuffled together, and the same tree cannot reach 2012 demand when trained on 2011 alone, because a terminal region can only average values it has already seen. That is a property of the model class rather than a tuning failure, and it explains why the simplest candidate transferred best.

6.3 What can be improved

The strongest limitations are evidential. More years, archived weather forecasts and repeated forward windows would test whether the temporal result reflects this transition or a permanent property. Station-level planning would additionally need identifiers, dock capacity, trip origins and destinations, transit connections and local demand, which is why the ITDP and NACTO guidance points beyond this dataset. Category-4 weather remains outside the supported domain.

Methodologically, trend-seasonality decomposition, differenced targets and forecast-aware inputs are the natural next comparisons.

The extension I would most like to build is a small service that turns the frozen primary model into something usable: send it a date, its season and a weather forecast, and receive an estimated daily volume with the supported-domain guard applied. Appendix C sketches that idea, including the parts this study shows it must not claim.

Appendix A - Glossary

Term	Meaning in this report
MAE	Mean absolute error: average miss in rentals per day.
MSE / RMSE	Squared error and its root; large misses receive greater weight.
R-squared	Variance explained relative to the evaluated sample mean; it may be negative.
Holdout	Data isolated from model, feature and hyperparameter selection and used once for final scoring.
Interpolation	Prediction within a domain represented in training.
Extrapolation	Prediction beyond the observed training range or period.
Autoregressive feature	Input constructed from demand observed before the target day.
TimeSeriesSplit	Expanding-window cross-validation that preserves temporal order.
Permutation importance	Model-agnostic holdout score degradation after shuffling an input.
SHAP	Additive prediction attribution used here only for a frozen tree winner.

Appendix B - Reproducibility and integrity evidence


```
In [23]: import sys
import sklearn
import matplotlib

assert len(primary_df) == 731
assert len(temporal_df) == 724
assert selection_indices.isdisjoint(set(primary_holdout_idx))
assert temporal_train.yr.eq(0).all()
assert temporal_train.dteday.max() < temporal_test.dteday.min()
assert all("casual" not in cols and "registered" not in cols
          for cols in list(PRIMARY_FEATURE_SETS.values()) + list(TEMPORAL_FEATURE_SETS.values()))
assert all(c not in cols for cols in PRIMARY_FEATURE_SETS.values() for c in AUTOREGRESSIVE_FEATURES)
assert cross_file.difference.eq(0).all() and len(cross_file) == 731
assert len(primary_holdout_table) == 5 and len(temporal_holdout_table) == 8

integrity = pd.DataFrame([
    ("Primary rows retained", 731, len(primary_df), "Pass"),
    ("Temporal rows after causal warm-up", 724, len(temporal_df), "Pass"),
    ("Primary holdout rows excluded from selection", 183, len(primary_holdout_idx), "Pass"),
    ("Temporal selection years", "2011 only", str(temporal_train.dteday.dt.year.unique().tolist()), "Pass"),
    ("Hourly/daily dates reconciled", 731, int(cross_file.difference.eq(0).sum()), "Pass"),
    ("Dates with incomplete hourly panels", 76, incomplete_dates, "Pass"),
    ("Omitted zero-demand hourly rows", 165, omitted_hour_rows, "Pass"),
    ("Leakage components in predictor lists", 0, 0, "Pass"),
], columns=["check", "expected", "observed", "result"])
integrity.to_csv(OUTPUT_DIR / "integrity_checks_v4.csv", index=False)
display(integrity)
print("Python:", sys.version.split()[0], "| pandas:", pd.__version__,
      "| scikit-learn:", sklearn.__version__, "| matplotlib:", matplotlib.__version__)
print("RANDOM_STATE:", RANDOM_STATE)
```

	check	expected	observed	result
0	Primary rows retained	731	731	Pass
1	Temporal rows after causal warm-up	724	724	Pass
2	Primary holdout rows excluded from selection	183	183	Pass
3	Temporal selection years	2011 only	[2011]	Pass
4	Hourly/daily dates reconciled	731	731	Pass
5	Dates with incomplete hourly panels	76	76	Pass
6	Omitted zero-demand hourly rows	165	165	Pass
7	Leakage components in predictor lists	0	0	Pass

Python: 3.14.3 | pandas: 3.0.3 | scikit-learn: 1.9.0 | matplotlib: 3.11.0
RANDOM_STATE: 42

Appendix C - Proposed extension: a demand estimation service

This appendix sketches the engineering extension named in Section 6.3. It is a proposal, not a deployed system, and the evidence in Section 5 constrains what it may offer.

C.1 What it would do

A caller sends a date, its season and a weather forecast; the service returns an estimated daily rental volume from the frozen primary model, together with the supported input domain.

```
POST /estimate
{ "date": "2013-04-18", "season": 2, "weathersit": 1,
  "temp": 0.58, "atemp": 0.55, "hum": 0.61, "windspeed": 0.18 }

-> { "estimated_daily_rentals": 5120,
     "expected_absolute_error": 434,
     "basis": "conditional estimate, observed 2011-2012 distribution",
     "supported_weathersit": [1, 2, 3] }
```

Returning the expected error alongside the point estimate matters more than the estimate itself. A planner who receives 5,120 without 434 will read a precision the model does not have.

C.2 Proposed layers

Layer	Technology	Purpose
Model artefact	scikit-learn pipeline, joblib	The frozen Gradient Boosting configuration and its preprocessing, serialised together
Contract	JSON	Feature order, supported <code>weathersit</code> categories, holdout metrics, source notebook hash
API	FastAPI, Pydantic	Typed request validation and the domain guard
Interface	Streamlit	A form for planners who will not call an API
Monitoring	Scheduled job	Daily scoring of yesterday's estimate against realised demand

The pattern follows my Sommelier API project from Assessments 1 and 2, where the same separation between a framework-agnostic model core and thin serving surfaces is already implemented.

C.3 What this study forbids it from claiming

The temporal check is the reason this appendix is short about forecasting.

- The service may answer **conditional** questions: given these conditions, what does the 2011-2012 relationship imply? That is the question the primary experiment evaluated.

- It may **not** be presented as a next-day forecaster. Section 4.2 measured the frozen model at 23.5% worse than a rolling seven-day mean on 2012. A forecasting endpoint would either serve the rolling mean or wait for new forward evidence.
- Any request outside `weathersit` 1 to 3 must be rejected for human review rather than scored, because category 4 never appears in training.

C.4 Geolocation and alerts

Two ideas worth recording, with their preconditions.

Station-level geolocation. Plotting expected demand onto a map through a mapping API is straightforward as an interface, and unsupported as an inference. This dataset carries no station identifier, coordinate, dock capacity or trip origin, so nothing in this notebook can attribute system demand to a location. Delivering it honestly means acquiring station-level data first; ITDP and NACTO describe the network, spacing and capacity variables that work would need (ITDP, 2018; NACTO, 2016).

Custom alerts. Two alert types are supported by evidence already collected. A *domain alert* fires when a request falls outside the supported weather categories. A *drift alert* fires when the rolling error of the deployed estimate exceeds its holdout MAE over a defined window, which is the mechanism that would have caught the temporal failure in production rather than in a notebook. A demand-threshold alert for staffing would be useful but needs a stakeholder-defined threshold that this project does not have.

Academic Integrity Declaration

I declare that this submission is my own work. All sources of information, ideas and code have been acknowledged. The dataset is publicly available from the UCI Machine Learning Repository. The analysis, modelling decisions, interpretation and written commentary are my own.

Statement of Acknowledgement

Analysis was performed in Python with pandas, NumPy, scikit-learn, Matplotlib, Seaborn and SHAP.

I acknowledge that I have used the following AI tool(s) in the creation of this report:

- Anthropic Claude Opus 5
- OpenAI ChatGPT Codex 5.6

Both tools were used to assist with understanding ML concepts, challenging the experimental design, auditing empirical claims against executed outputs, structuring the technical pipeline, improving clarity of academic language, and supporting APA 7th referencing conventions.

Prompt examples:

1. "Explain why a random split over a two-year time series supports a conditional demand estimate but not a next-day forecast, and how to run both experiments so that neither contaminates the other."
2. "Show why a tree ensemble trained on 2011 cannot predict 2012 demand above its training maximum, and propose a way to measure that ceiling instead of asserting it."
3. "Build a causal seven-day rolling baseline for a day-ahead comparison using shift before rolling, and write the assertion that proves the window never includes the target day."

I confirm that the use of these tools has been in accordance with the Torrens University Australia Academic Integrity Policy and TUA, Think and MDS's Position Paper on the Use of AI. I confirm that the final output is authored by me and represents my own critical thinking, analysis, and synthesis of sources. I take full responsibility for the final content of this report.

References

Capital Bikeshare. (n.d.). *How it works*. <https://capitalbikeshare.com/how-it-works>

Chapman, P., Clinton, J., Kerber, R., Khabaza, T., Reinartz, T., Shearer, C., & Wirth, R. (2000). *CRISP-DM 1.0: Step-by-step data mining guide*. SPSS Inc.

Fanaee-T, H., & Gama, J. (2014). Event labeling combining ensemble detectors and background knowledge. *Progress in Artificial Intelligence*, 2, 113-127.
<https://doi.org/10.1007/s13748-013-0040-3>

Friedman, J. H. (2001). Greedy function approximation: A gradient boosting machine. *Annals of Statistics*, 29(5), 1189-1232. <https://doi.org/10.1214/aos/1013203451>

Institute for Transportation & Development Policy. (2018). *The bikeshare planning guide*.
<https://itdp.org/publication/the-bike-share-planning-guide/>

Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30, 4765-4774.

National Association of City Transportation Officials. (2016). *Bike share station siting guide*. <https://nacto.org/publication/bike-share-station-siting-guide/>

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825-2830.

University of California, Irvine. (n.d.). *Bike Sharing Dataset*. UCI Machine Learning Repository. <https://archive.ics.uci.edu/dataset/275/bike+sharing+dataset>